

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



deeplearning.ai

Case Studies

Why look at
case studies?

Outline

Classic networks:

- LeNet-5 ←
- AlexNet ←
- VGG ←

ResNet (152)

Inception




deeplearning.ai

Case Studies

Classic networks

LeNet - 5



60K parameters.

$n_H, n_W \downarrow$ $n_C \uparrow$


conv pool conv pool fc fc output

Advanced: sigmoid/tanh ReLU

II, III.

↓

AlexNet




- Similar to LeNet, but much bigger.

- ReLU

- Multiple GPUs.

- Local Response Normalization (LRN)



~60M parameters



VGG - 16

16 layers that has weight




VGG-19

CONV = 3x3 filter, s = 1, same

MAX-POOL = 2x2, s = 2



224x224x3



$n_h, n_w \downarrow$

$n_c \uparrow$

~138M

a lot of hyperparameter

VGG-19 → bigger but work similar with VGG-16




deeplearning.ai

Case Studies


Residual Networks (ResNets)

Residual block



residual blocks take shortcut instead of main path

"short cut" / skip connection



$$z^{[l+1]} = W^{[l+1]} a^{[l]} + b^{[l+1]}$$


$$a^{[l+1]} = g(z^{[l+1]})$$

$$z^{[l+2]} = W^{[l+2]} a^{[l+1]} + b^{[l+2]}$$

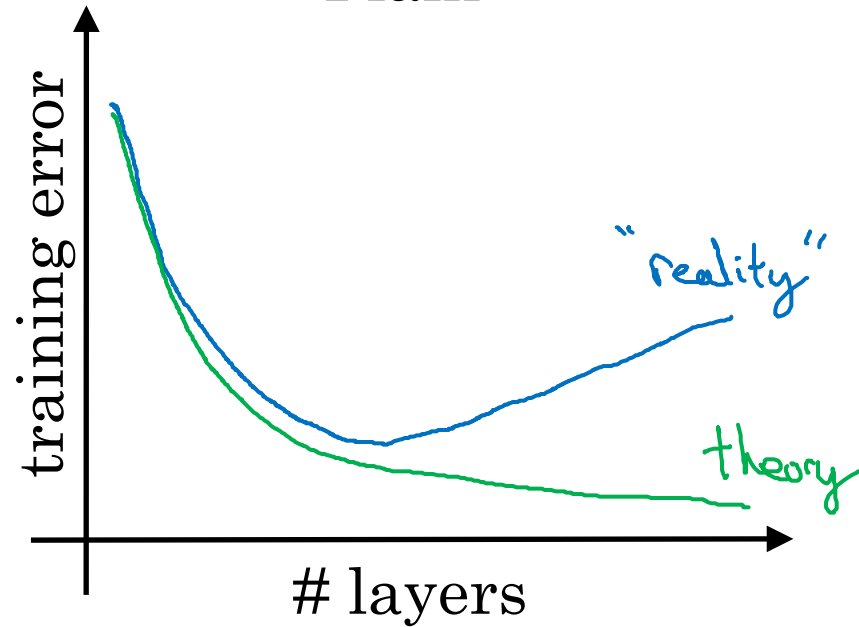
~~$$a^{[l+2]} = g(z^{[l+2]})$$~~

$$a^{[l+2]} = g(z^{[l+2]} + \underbrace{a^{[l]}}_{\text{skip}})$$

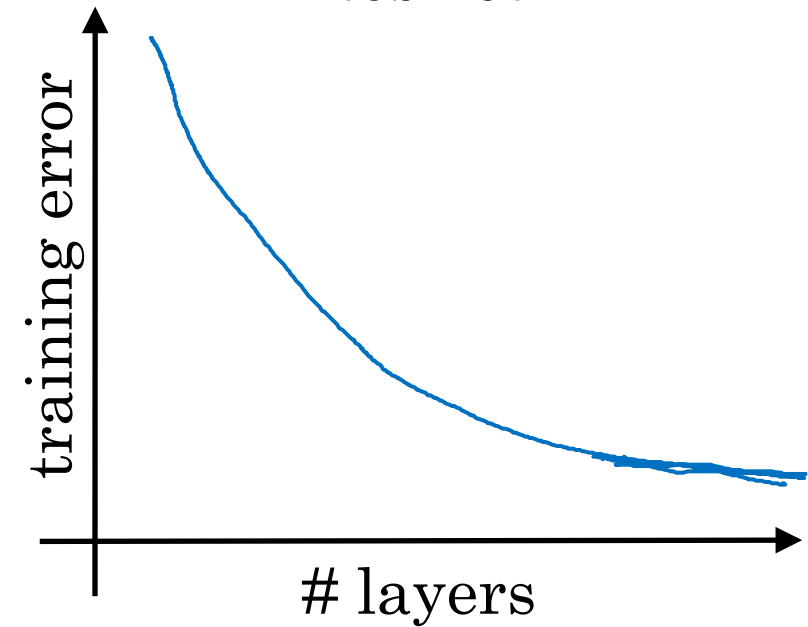
Residual Network



Plain



ResNet






deeplearning.ai

Case Studies

Why ResNets work

Why do residual networks work?



ReLU.

$a \geq 0$

$a^{[l+2]} = g(\underbrace{z^{[l+2]} + a^{[l]}}_{\text{same}})$

$= g(\underbrace{w^{[l+2]} + a^{[l+1]} + b^{[l+2]}}_{\text{If } w^{[l+2]}=0, b^{[l+2]}=0} + \underbrace{w_s^{[l+2]} a^{[l]}}_{\text{add } W_s \text{ if } a[l+1] \text{ is not the same d with } a[l]}) = g(a^{[l]})$

$= \underline{a^{[l]}}$

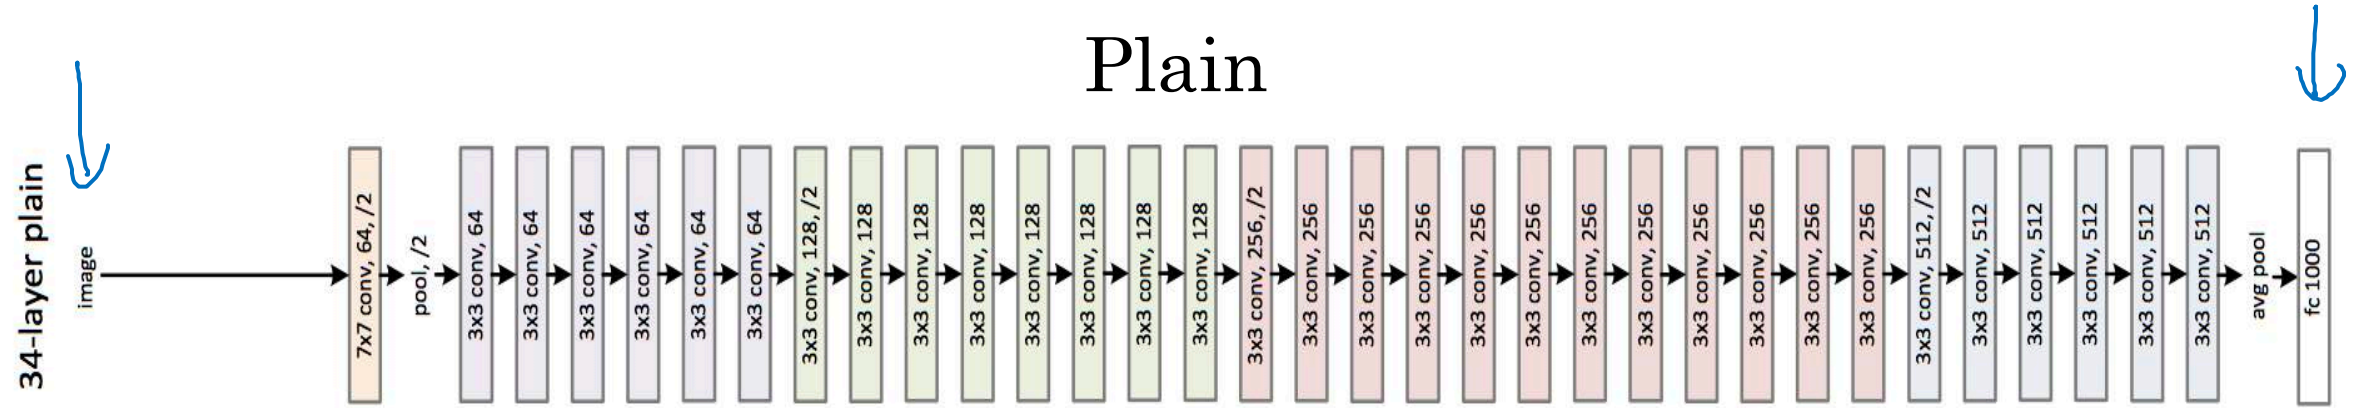
256

256 x 128

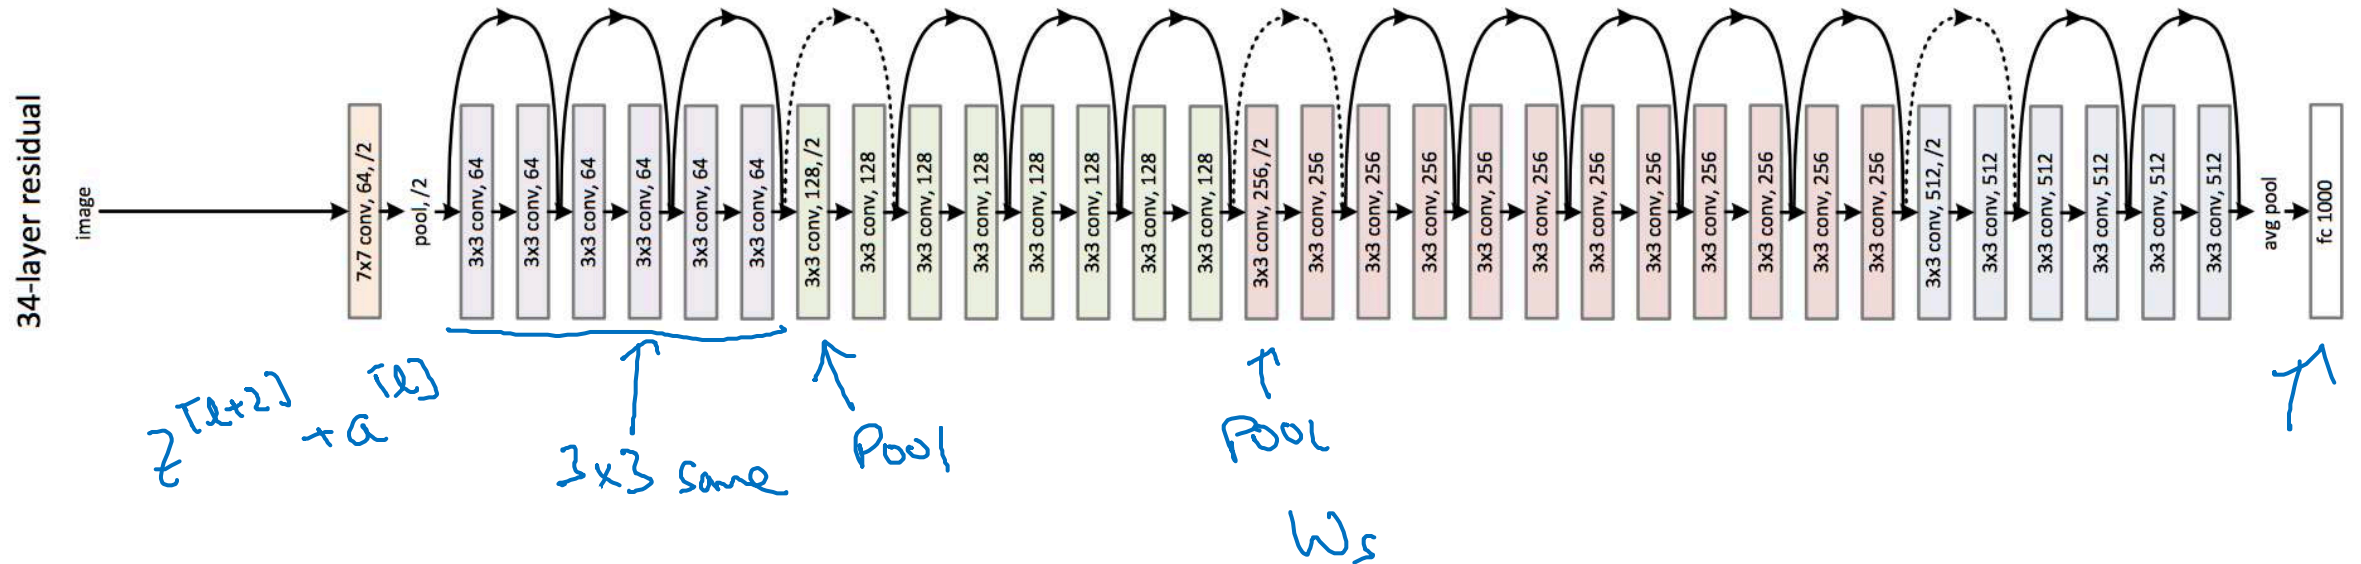
128

ResNet

Plain



ResNet





deeplearning.ai

Case Studies


Network in Network
and 1×1 convolutions

Why does a 1 × 1 convolution do?

1	2	3	6	5	8
3	5	5	1	3	4
2	1	3	4	9	3
4	7	8	5	7	9
1	5	3	7	4	8
5	4	9	8	3	5

$6 \times 6 \times 1$

like for each "bar"
input 32 value → output 1 value




$6 \times 6 \times 32$

*

2

=

32 → # filters.
 $n_c^{[l+1]}$




$1 \times 1 \times 32$

=

ReLU


Network in
Network

2	4	6	...		



$6 \times 6 \times \# \text{ filters}$

Using 1x1 convolutions



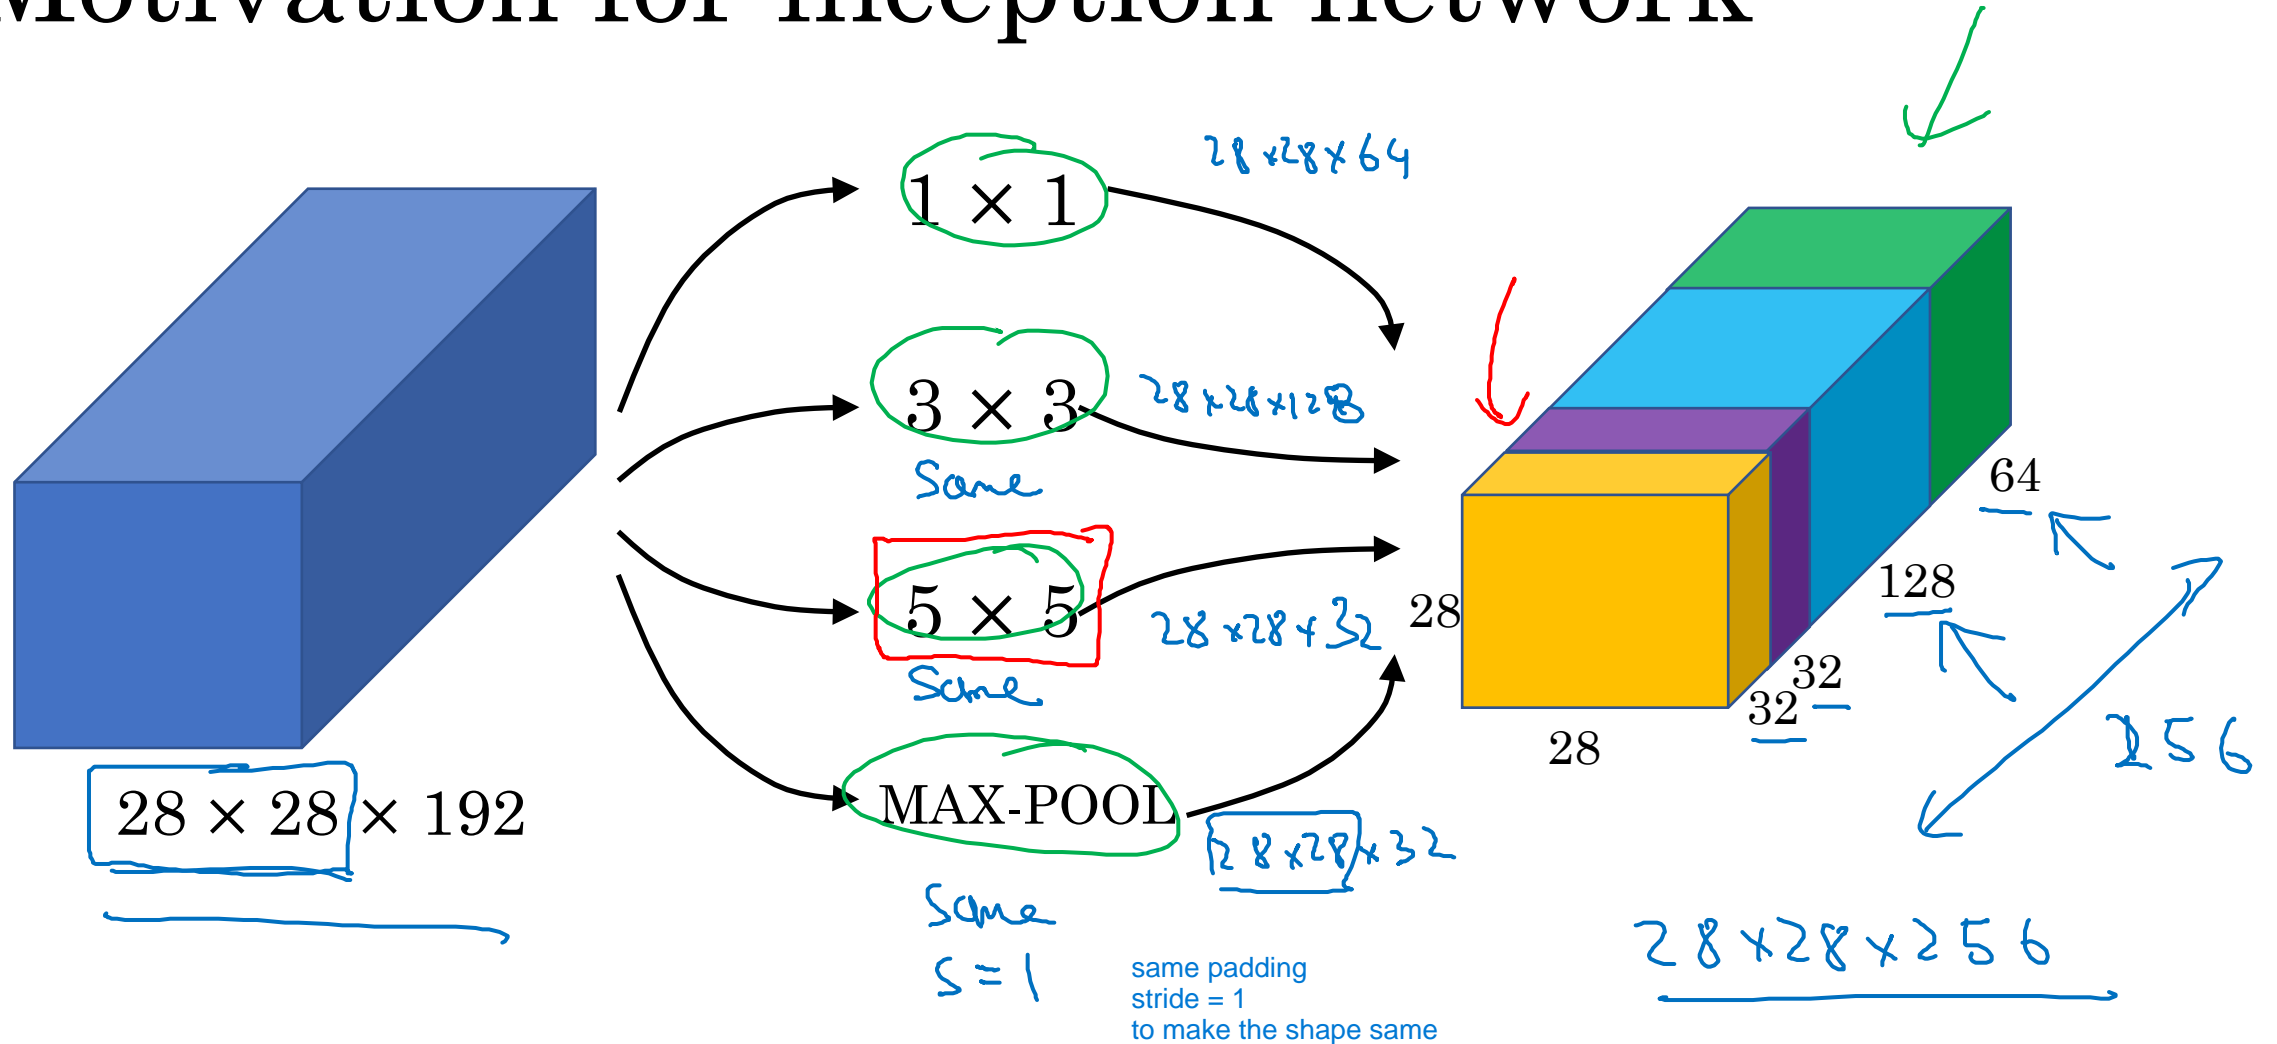


deeplearning.ai


Case Studies

Inception network motivation

Motivation for inception network




The problem of computational cost




Using 1×1 convolution

"bottleneck layer"




smaller intermediate layer
bottleneck layer




$$28 \times 28 \times 192$$

CONV
 1×1 ,
 $\rightarrow 16$,
 $\rightarrow 1 \times 1 \times 192$



$$28 \times 28 \times 16$$

CONV
 5×5 ,
32,
 $5 \times 5 \times 16$



$$28 \times 28 \times 32$$

$$28 \times 28 \times 16 \times 192 = 2.4M$$

$$28 \times 28 \times 32 \times 5 \times 5 \times 16 = 10.0M$$

$$12.4M$$

$$120M$$

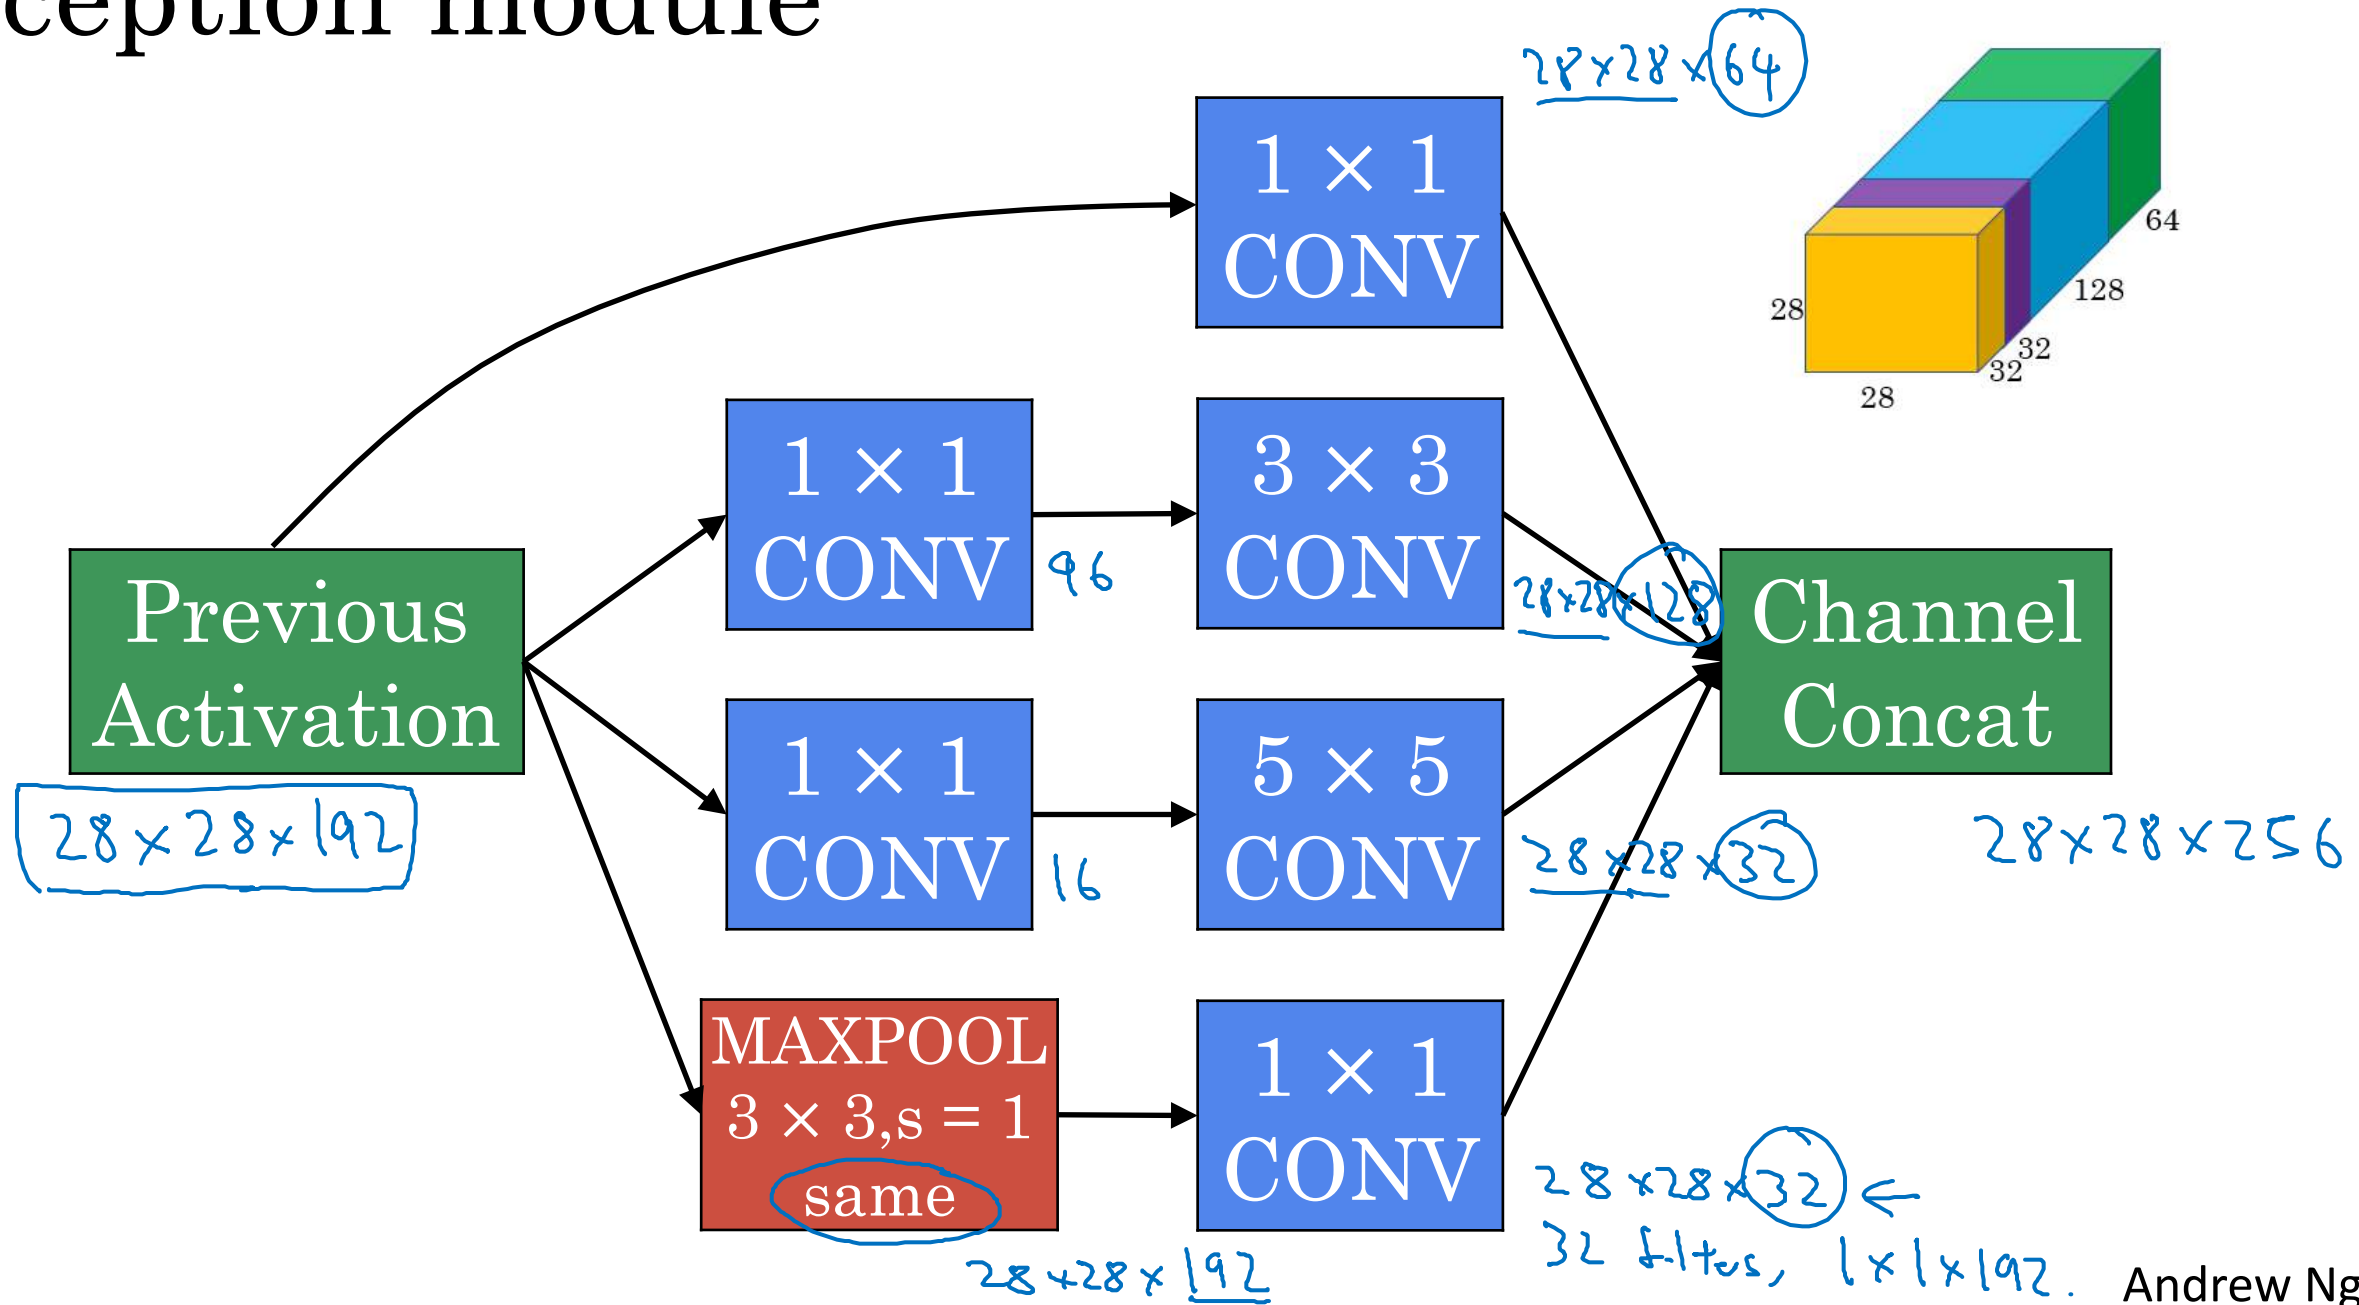


deeplearning.ai

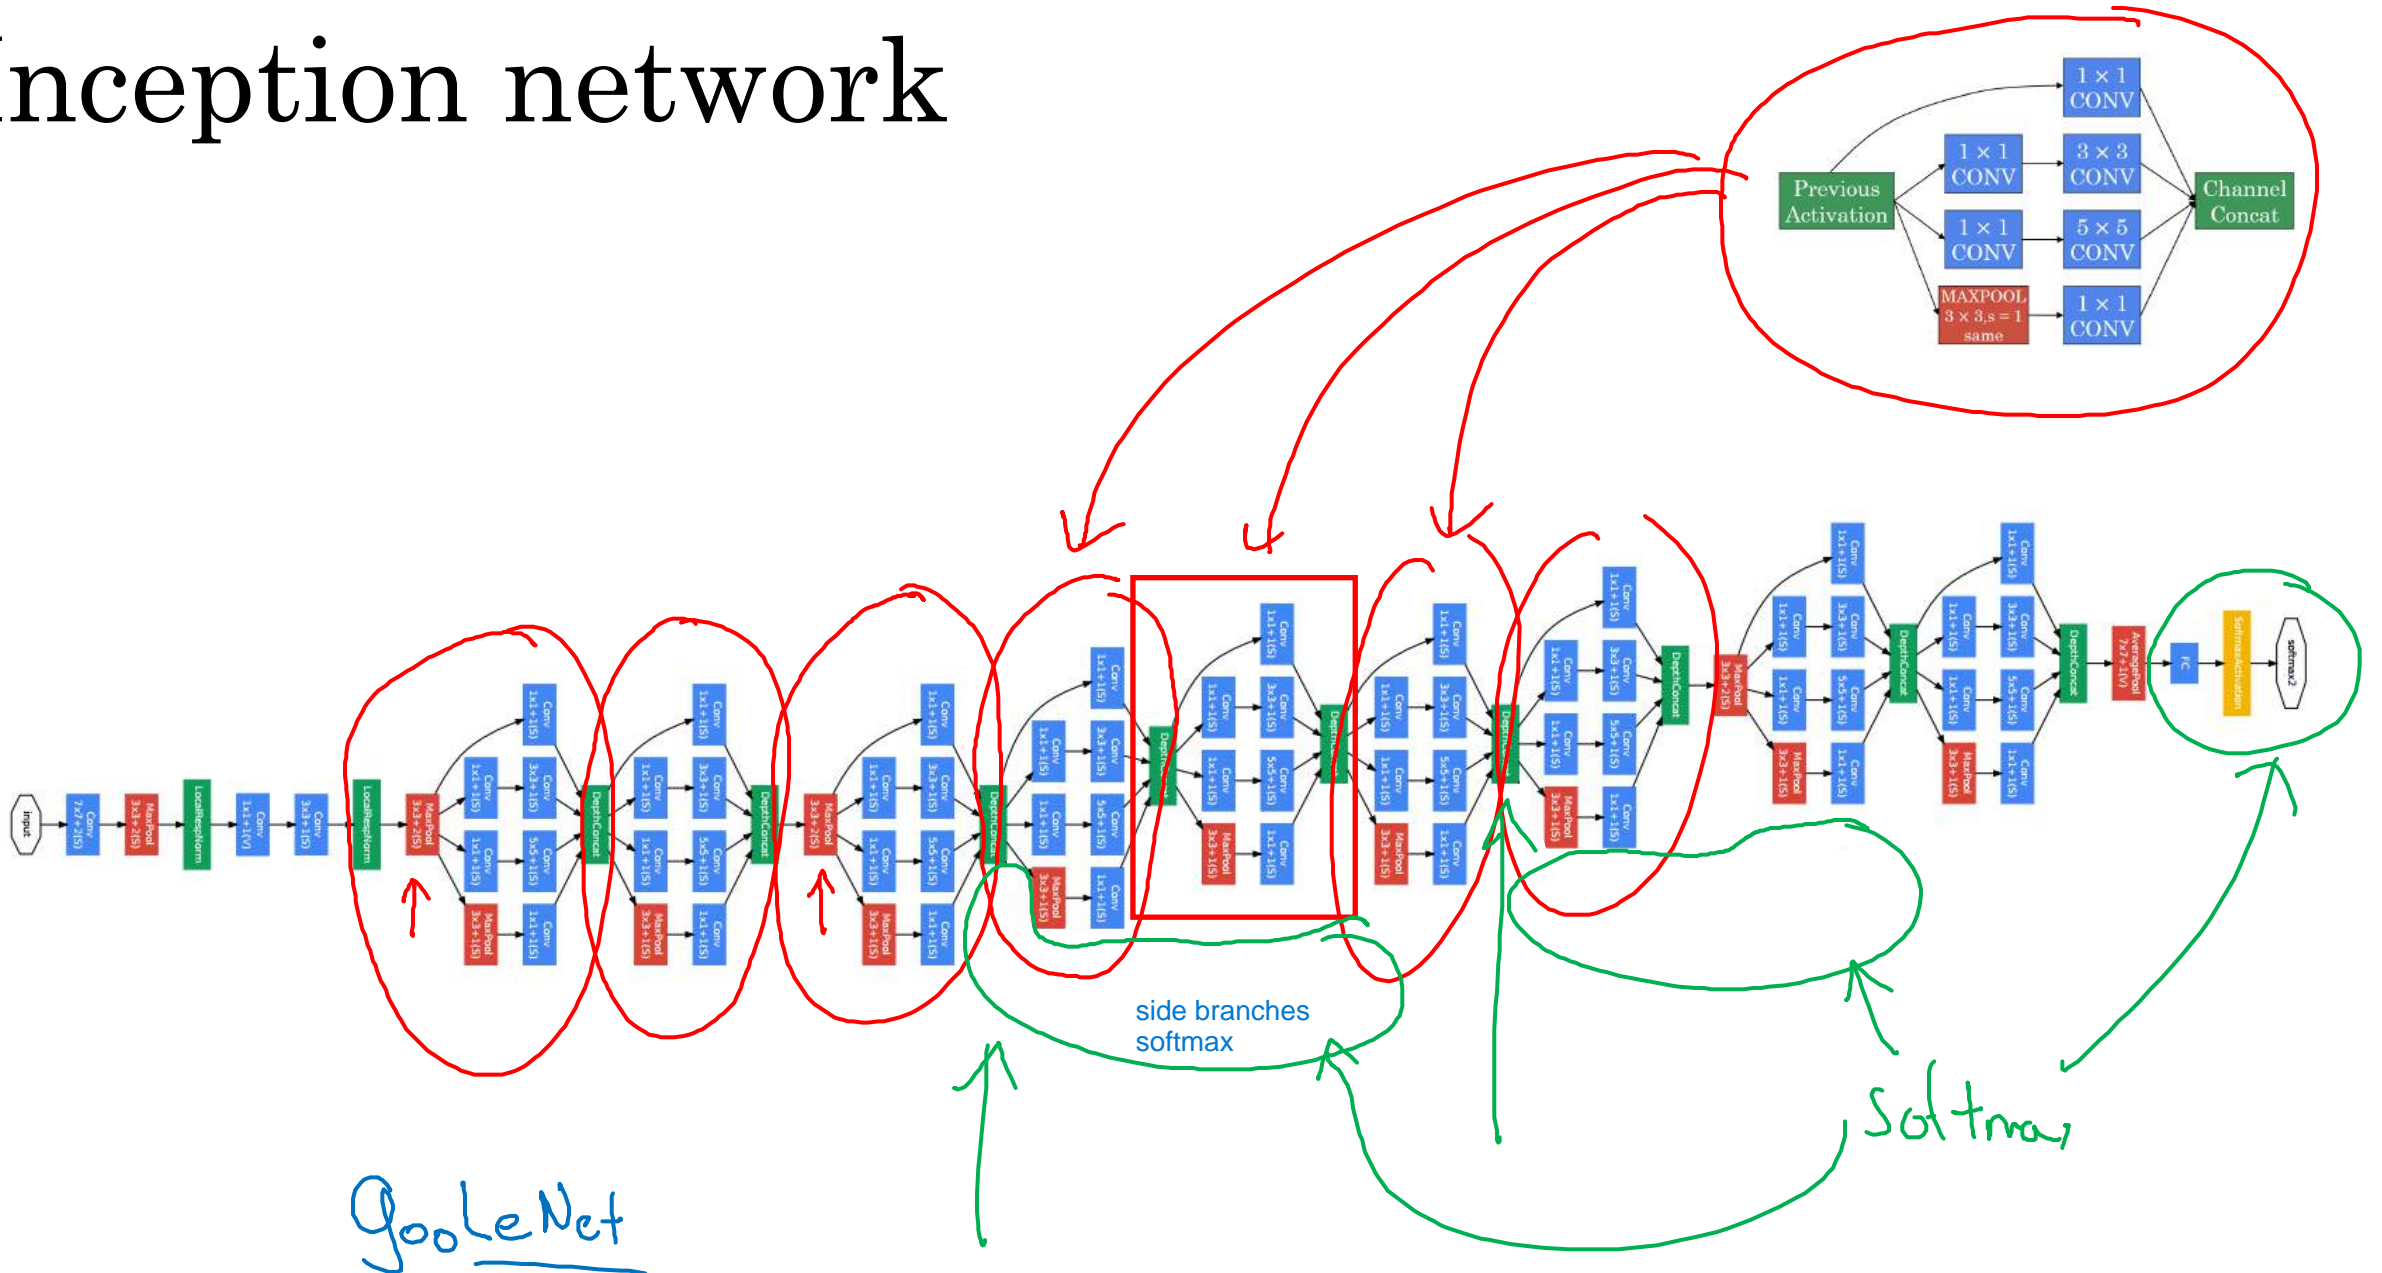
Case Studies

Inception network

Inception module



Inception network







deeplearning.ai

Convolutional Neural Networks

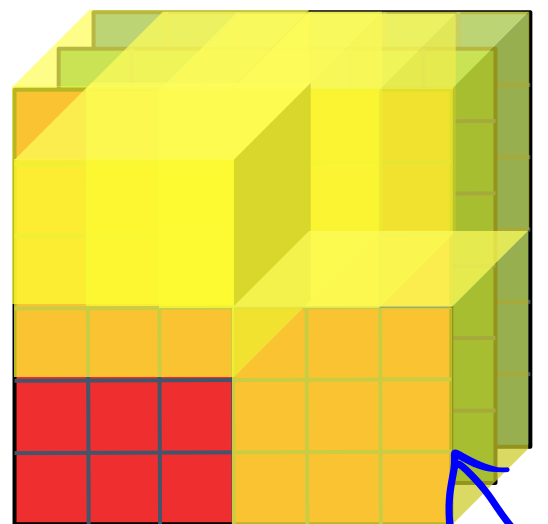
MobileNet

Motivation for MobileNets

- Low computational cost at deployment
- Useful for mobile and embedded vision applications
- Key idea: Normal vs. depthwise-separable convolutions




Normal Convolution




$6 \times 6 \times 3$
 $n \times n \times n_c$

*



$3 \times 3 \times 3$
 $f \times f \times n_c$

=



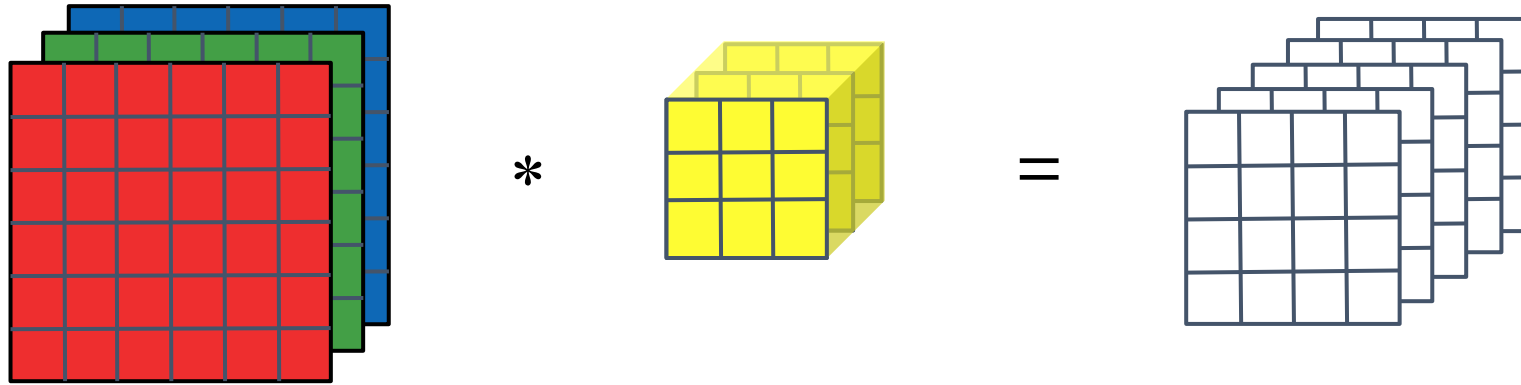
$4 \times 4 \times 5$
 $n_{out} \times n_{out} \times n_c'$

$$\begin{array}{ccccccc} \text{Computational cost} & = & \frac{\text{\#filter params}}{3 \times 3 \times 3} & \times & \frac{\text{\# filter positions}}{4 \times 4} & \times & \text{\# of filters} \\ \rightarrow \underline{2160} & = & \uparrow & & \uparrow & & \uparrow \\ & & & & & & 5 \end{array}$$

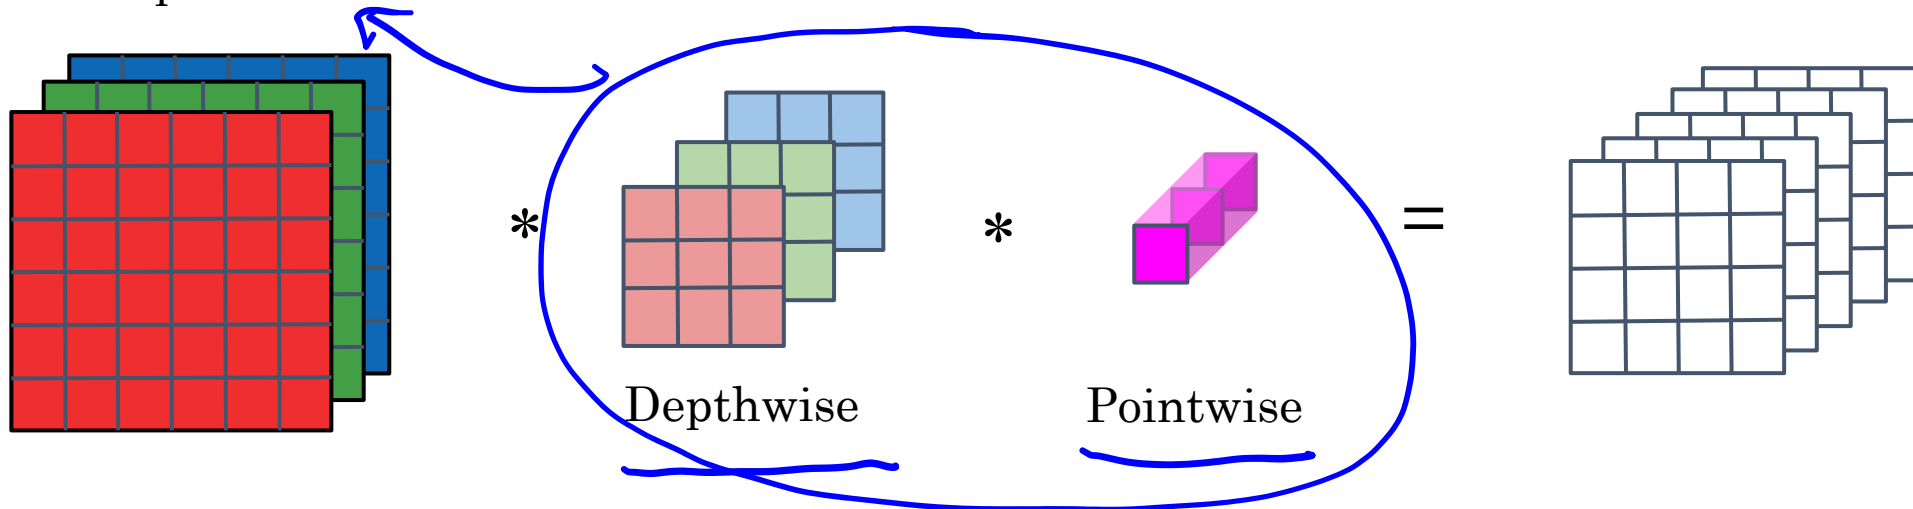
num of multiplication needed

Depthwise Separable Convolution

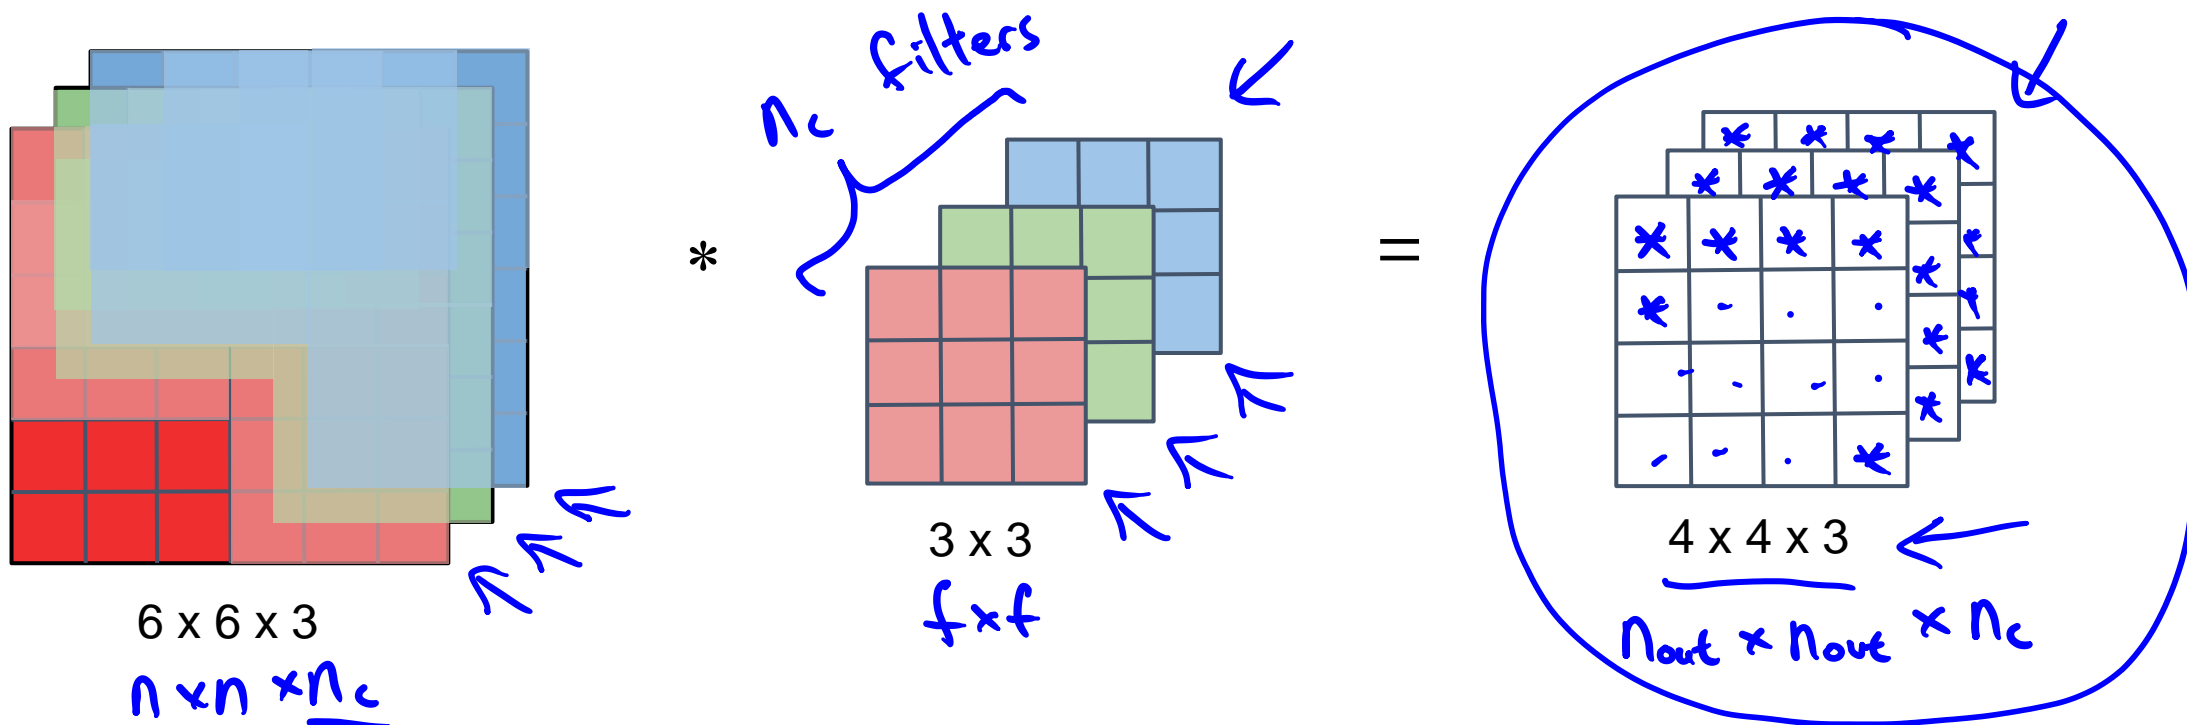
Normal Convolution



Depthwise Separable Convolution



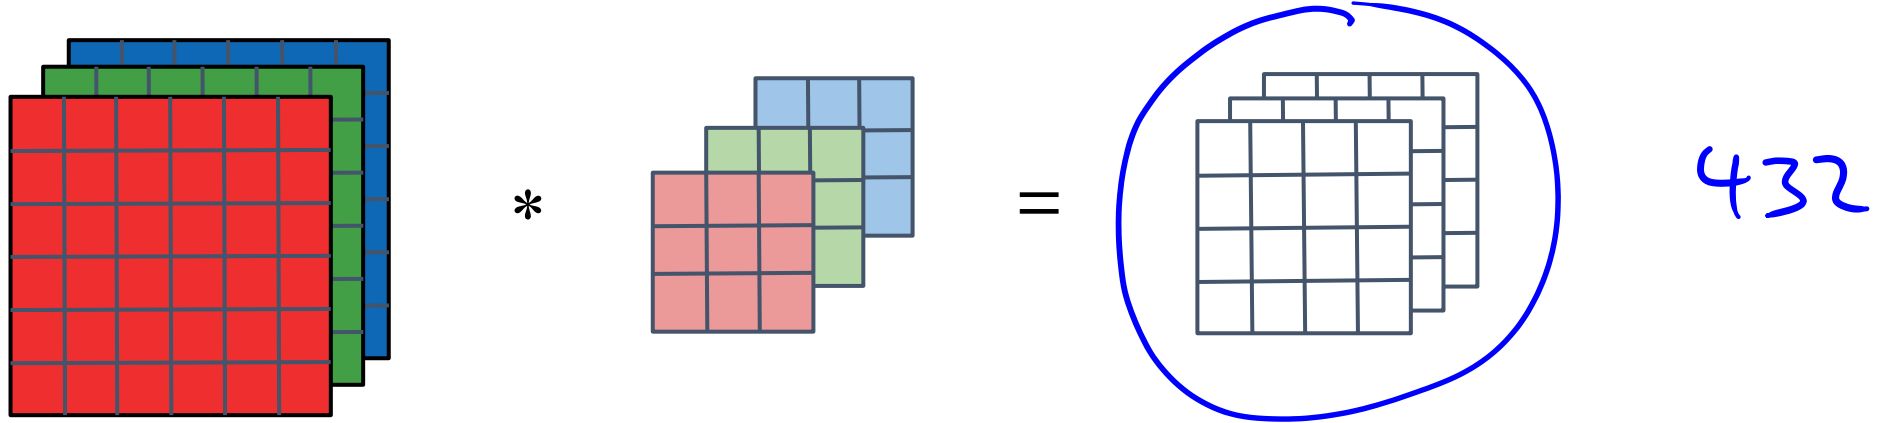
Depthwise Convolution




$$\begin{aligned}
 \text{Computational cost} &= \text{\#filter params} \times \text{\# filter positions} \times \text{\# of filters} \\
 432 &= \underbrace{3 \times 3} \times \underbrace{4 \times 4 \times 3}
 \end{aligned}$$

Depthwise Separable Convolution

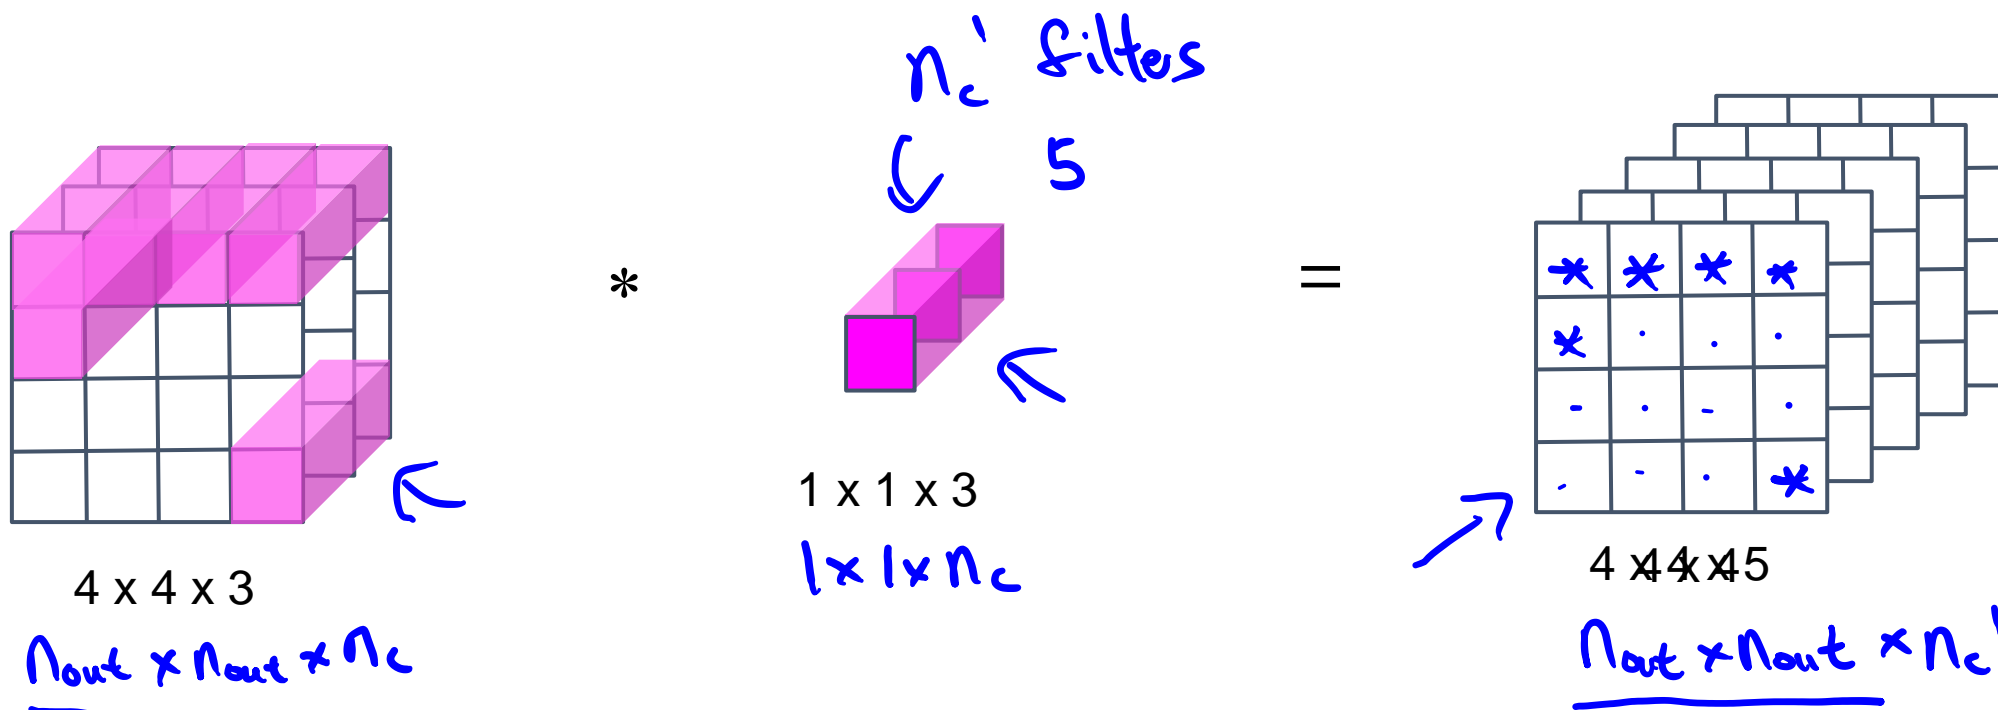
Depthwise Convolution



Pointwise Convolution



Pointwise Convolution

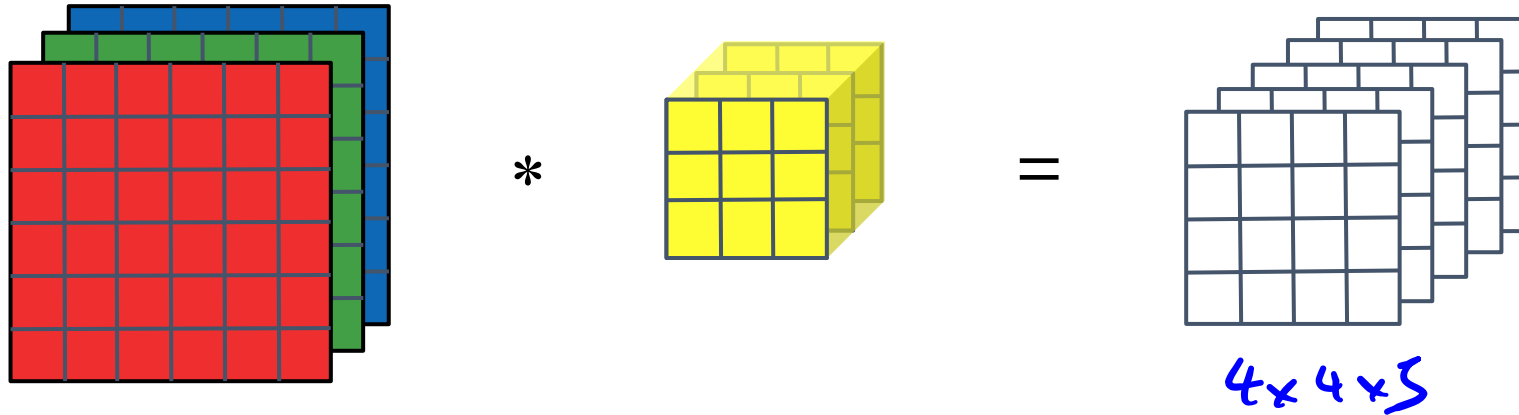


Computational cost = #filter params x # filter positions x # of filters

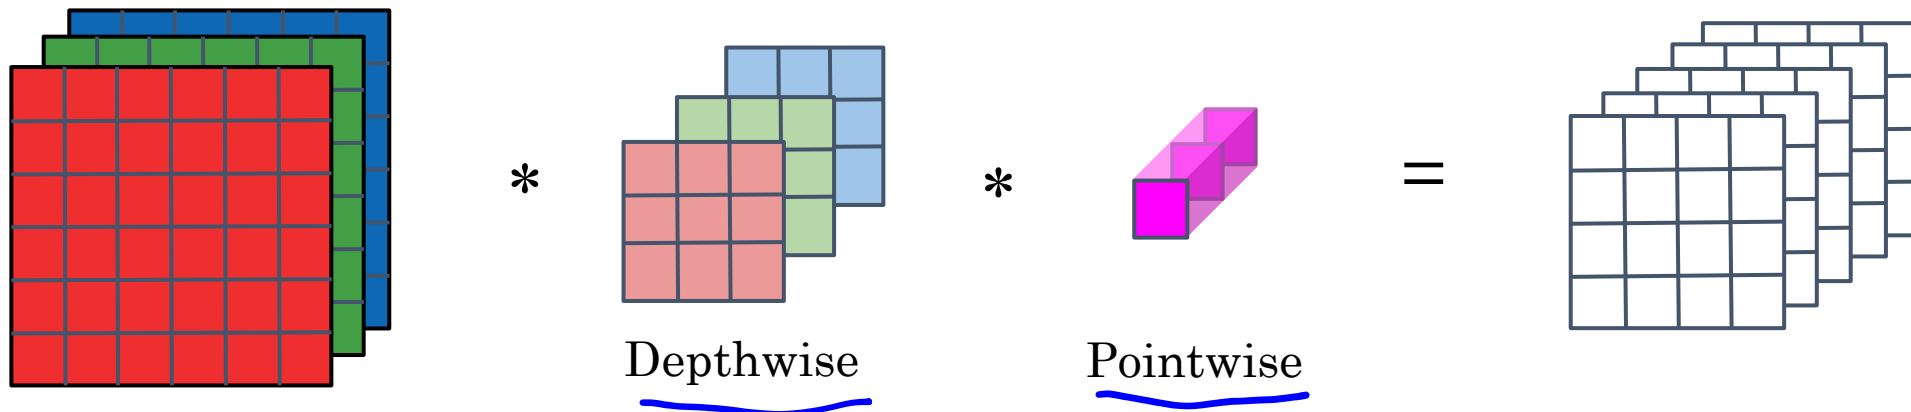
240 = $1 \times 1 \times 3$ x 4×4 x 5

Depthwise Separable Convolution

Normal Convolution



Depthwise Separable Convolution



Cost Summary

Cost of normal convolution

2160

Cost of depthwise separable convolution

depthwise + pointwise
432 + 240 = 672

$$\frac{672}{2160} = 0.31 \leftarrow$$


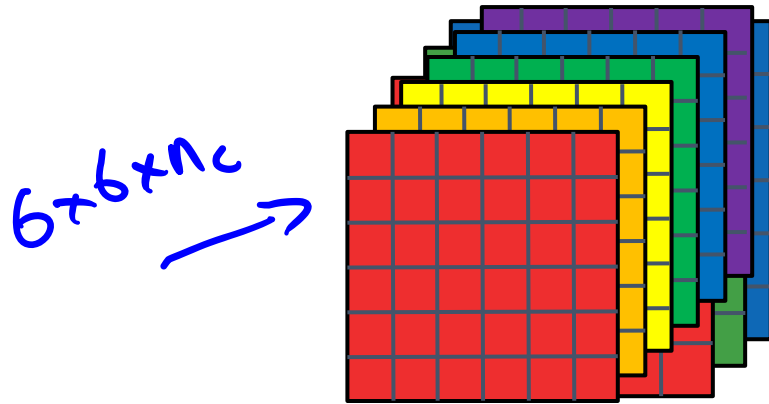
$$= \frac{1}{n_c} + \frac{1}{f^2}$$
$$\frac{1}{5} + \frac{1}{9}$$

$$= \frac{1}{512} + \frac{1}{3^2}$$


~10 times cheaper

Depthwise Separable Convolution



Depthwise Convolution




=



Pointwise Convolution



=






deeplearning.ai

Convolutional Neural Networks

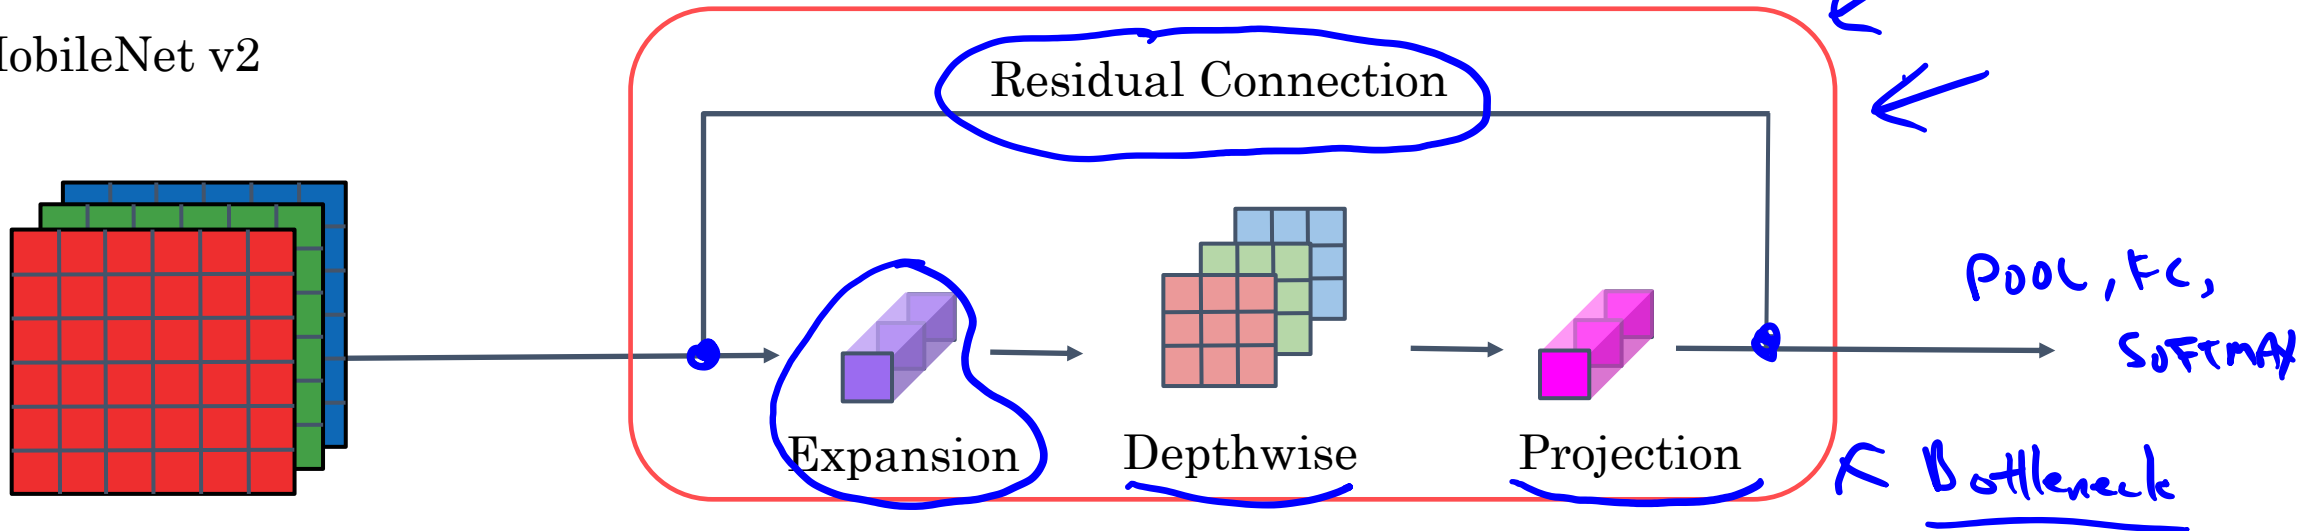
MobileNet Architecture

MobileNet

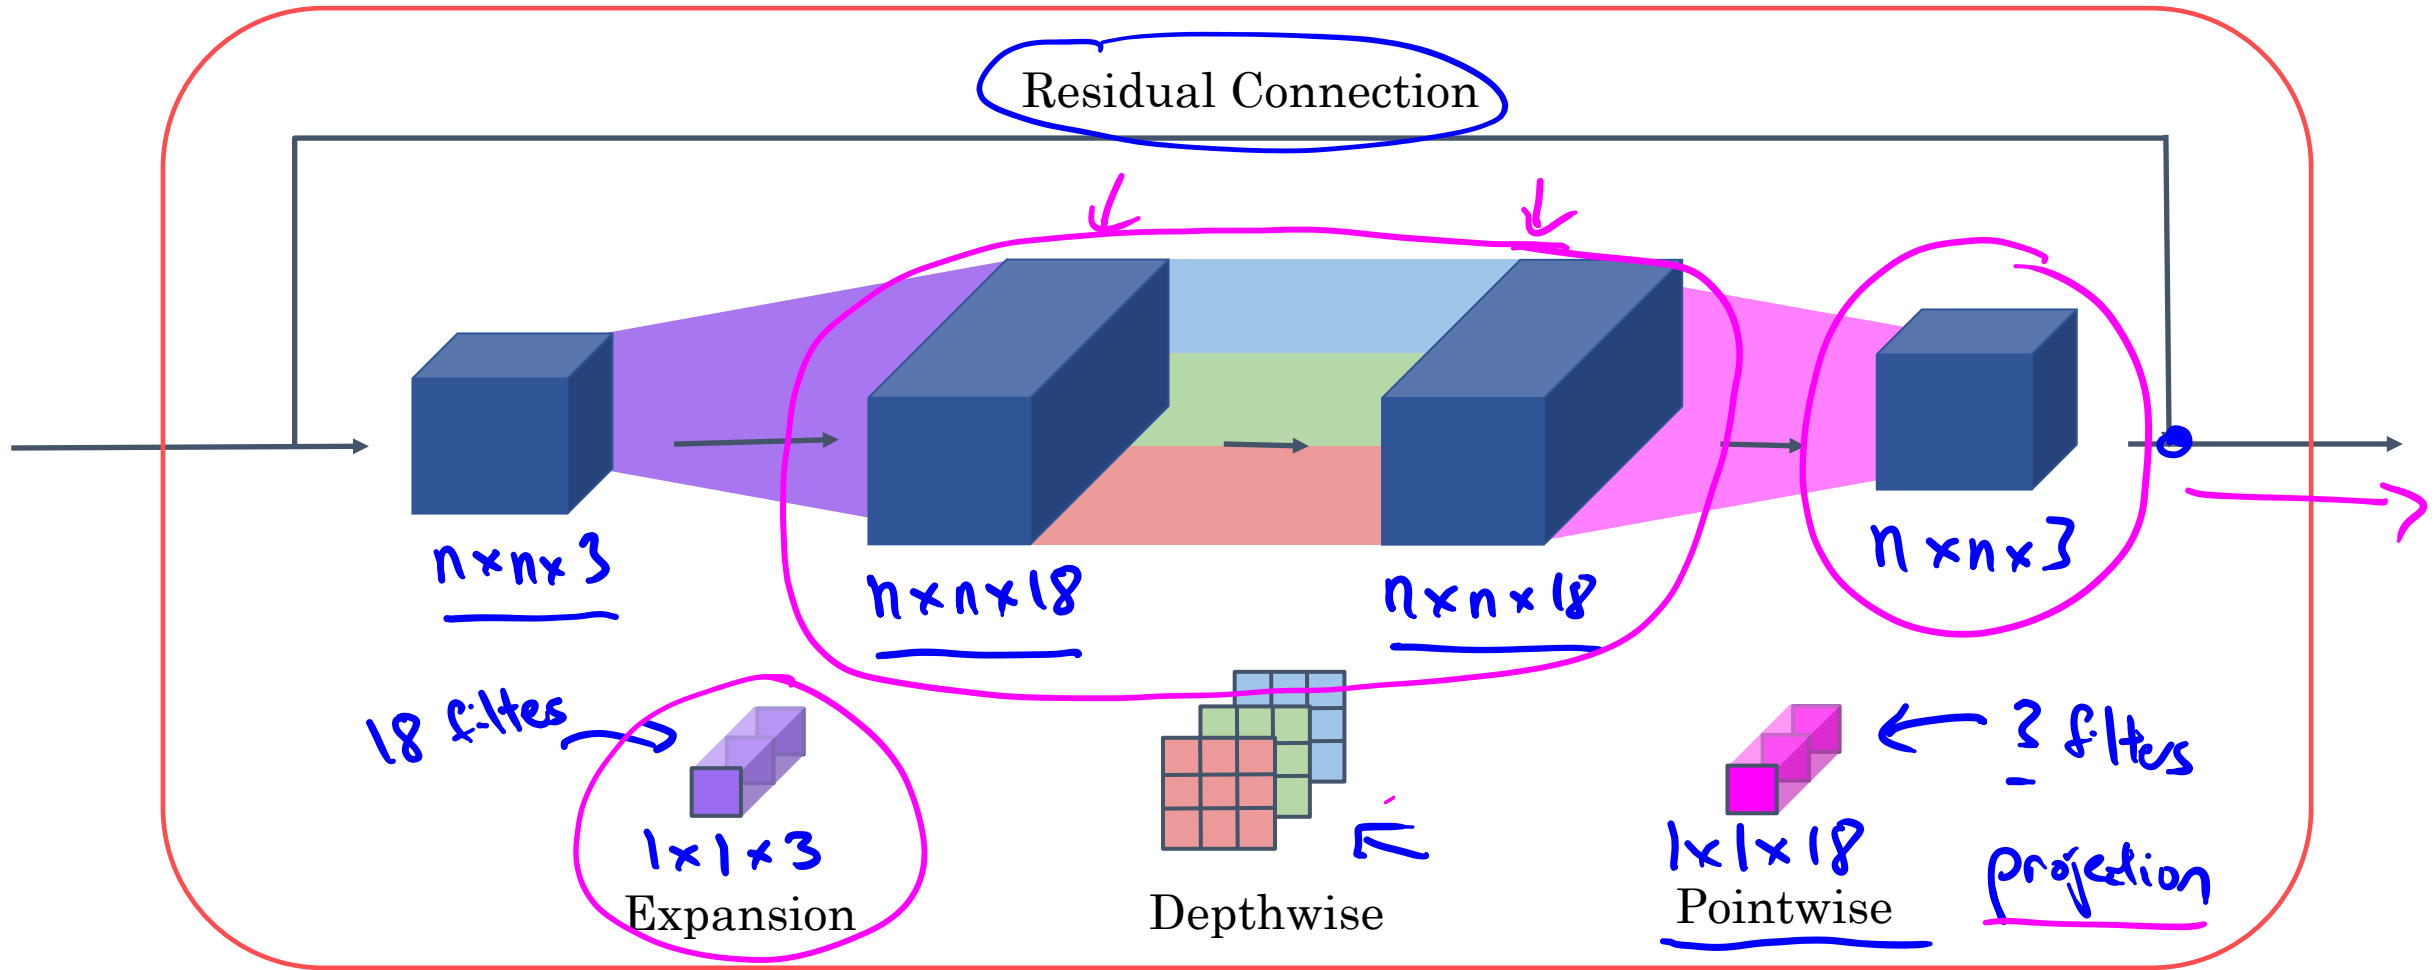
MobileNet v1



MobileNet v2

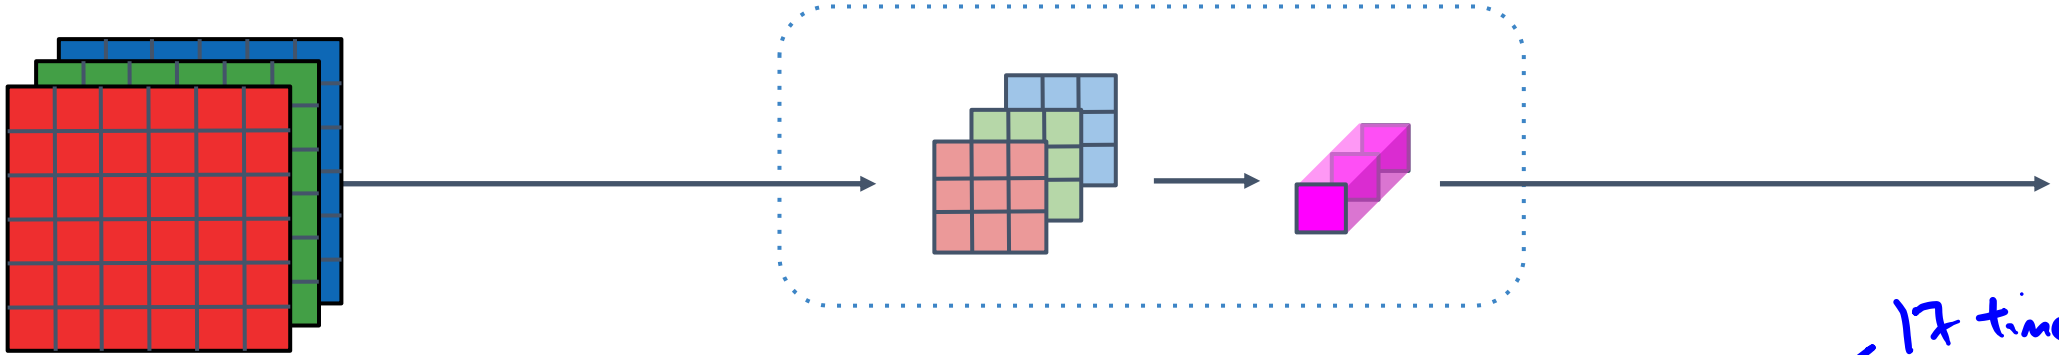


MobileNet v2 Bottleneck

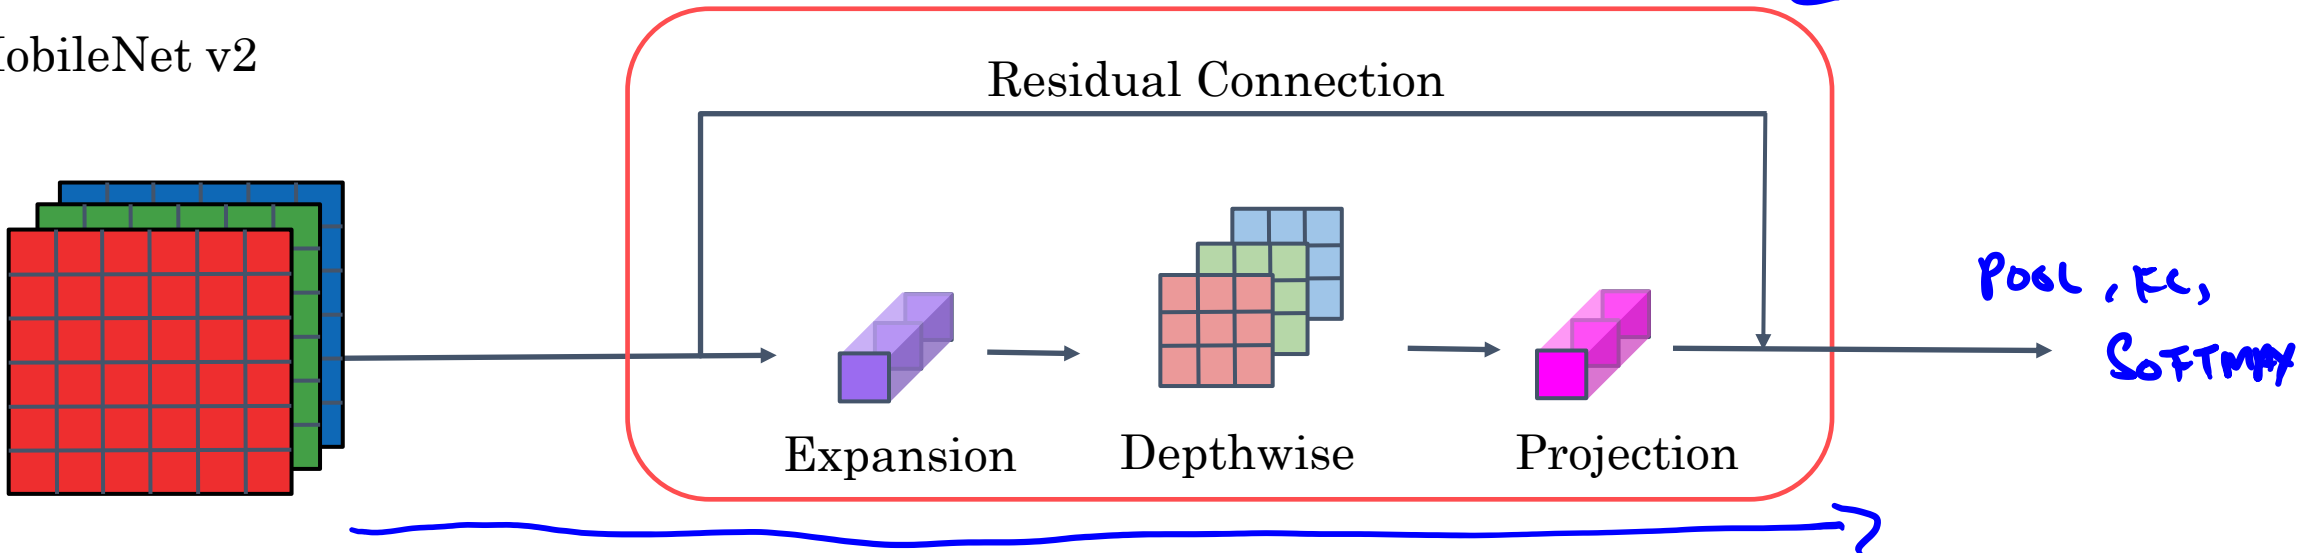


MobileNet

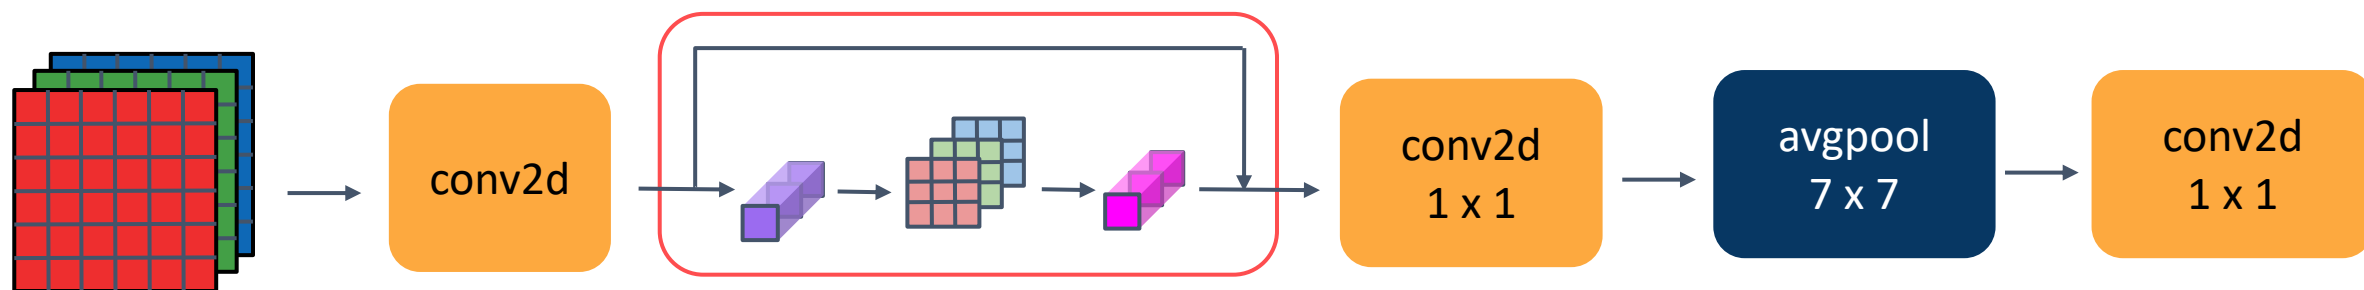
MobileNet v1



MobileNet v2



MobileNet v2 Full Architecture





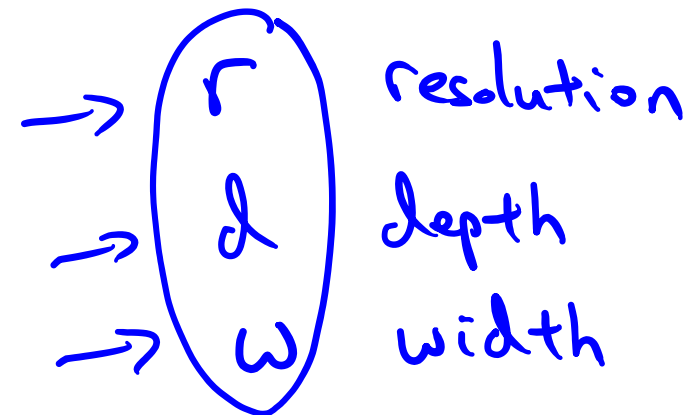

deeplearning.ai

Convolutional Neural Networks


EfficientNet

EfficientNet

Baseline



Compound Scaling





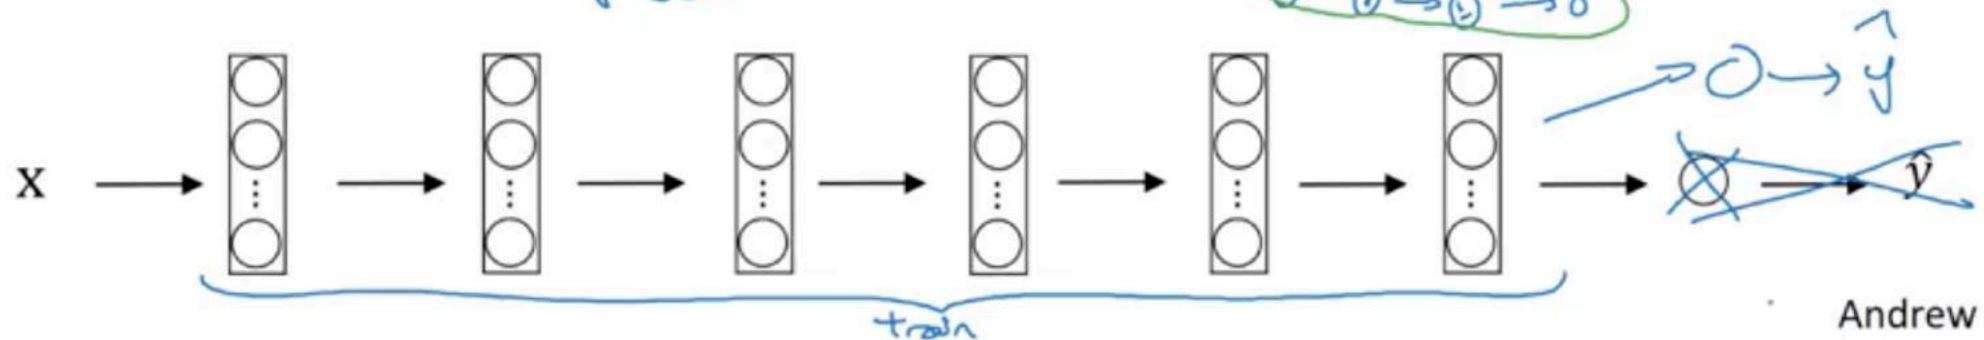
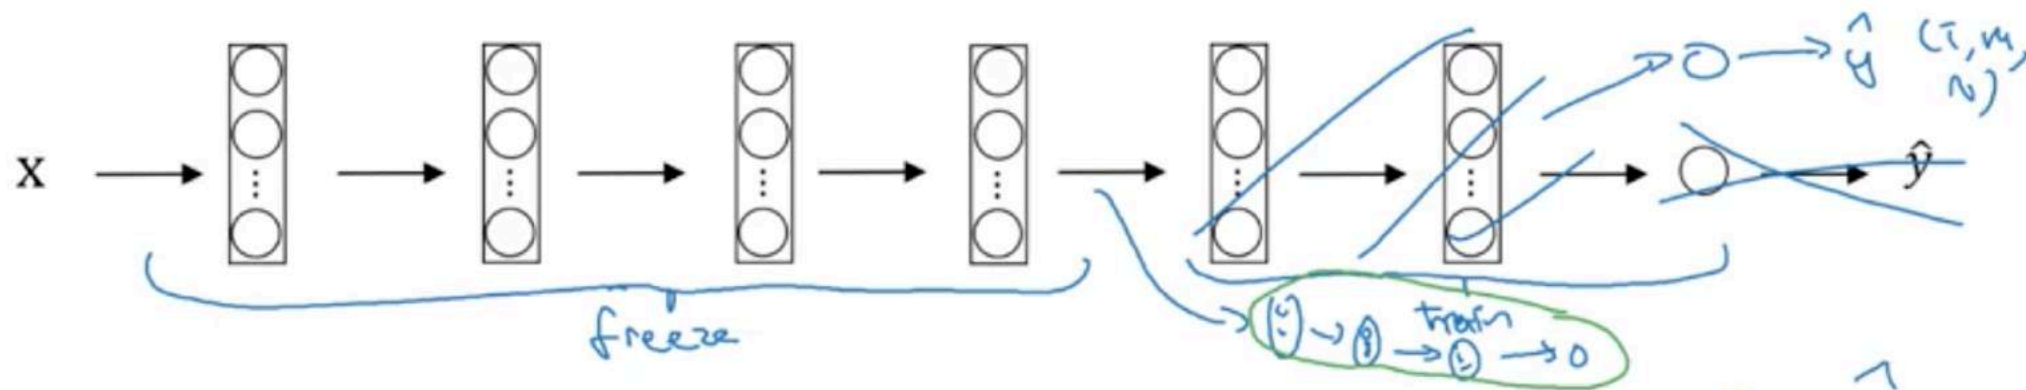
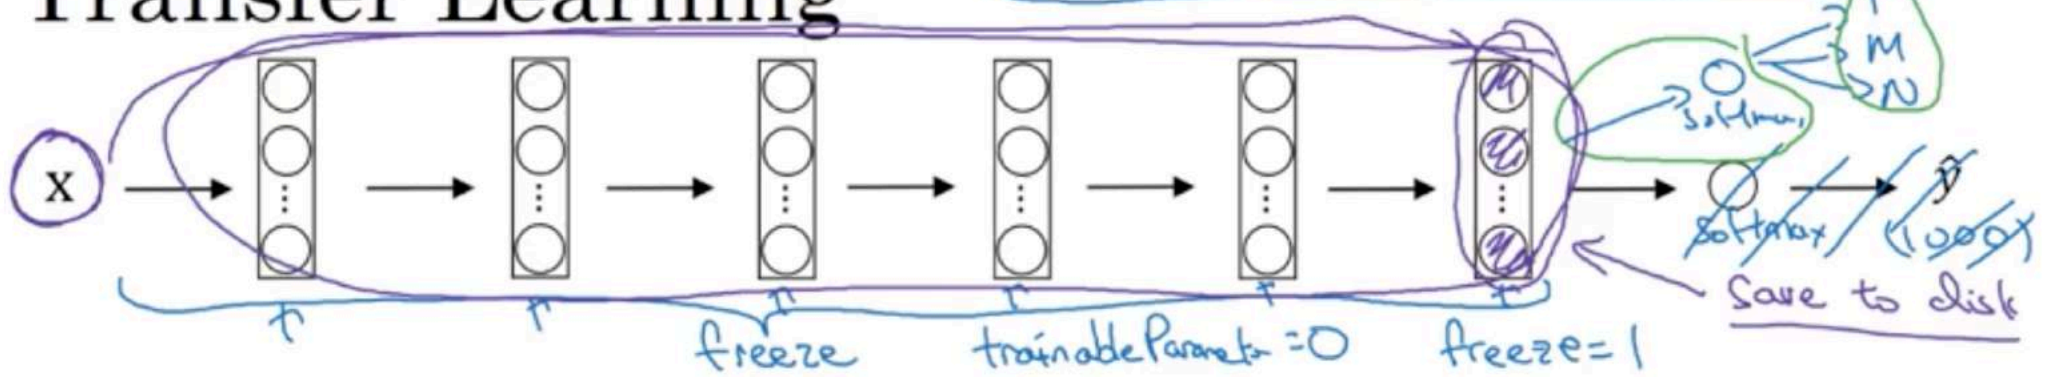
deeplearning.ai

Practical advice for using ConvNets

Transfer Learning

Transfer Learning

 Tigger
 Misty
Neither





deeplearning.ai

Practical advice for using ConvNets

Data augmentation

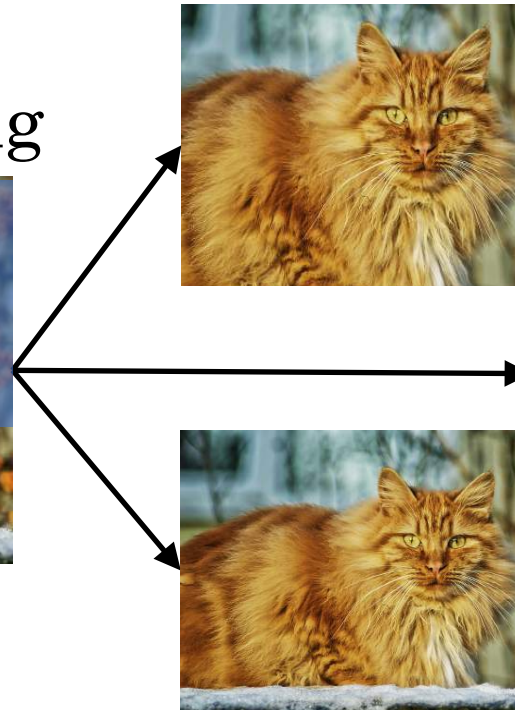
Common augmentation method

Mirroring



y

Random Cropping

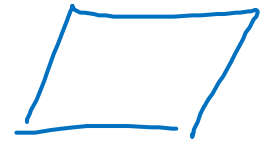


Rotation

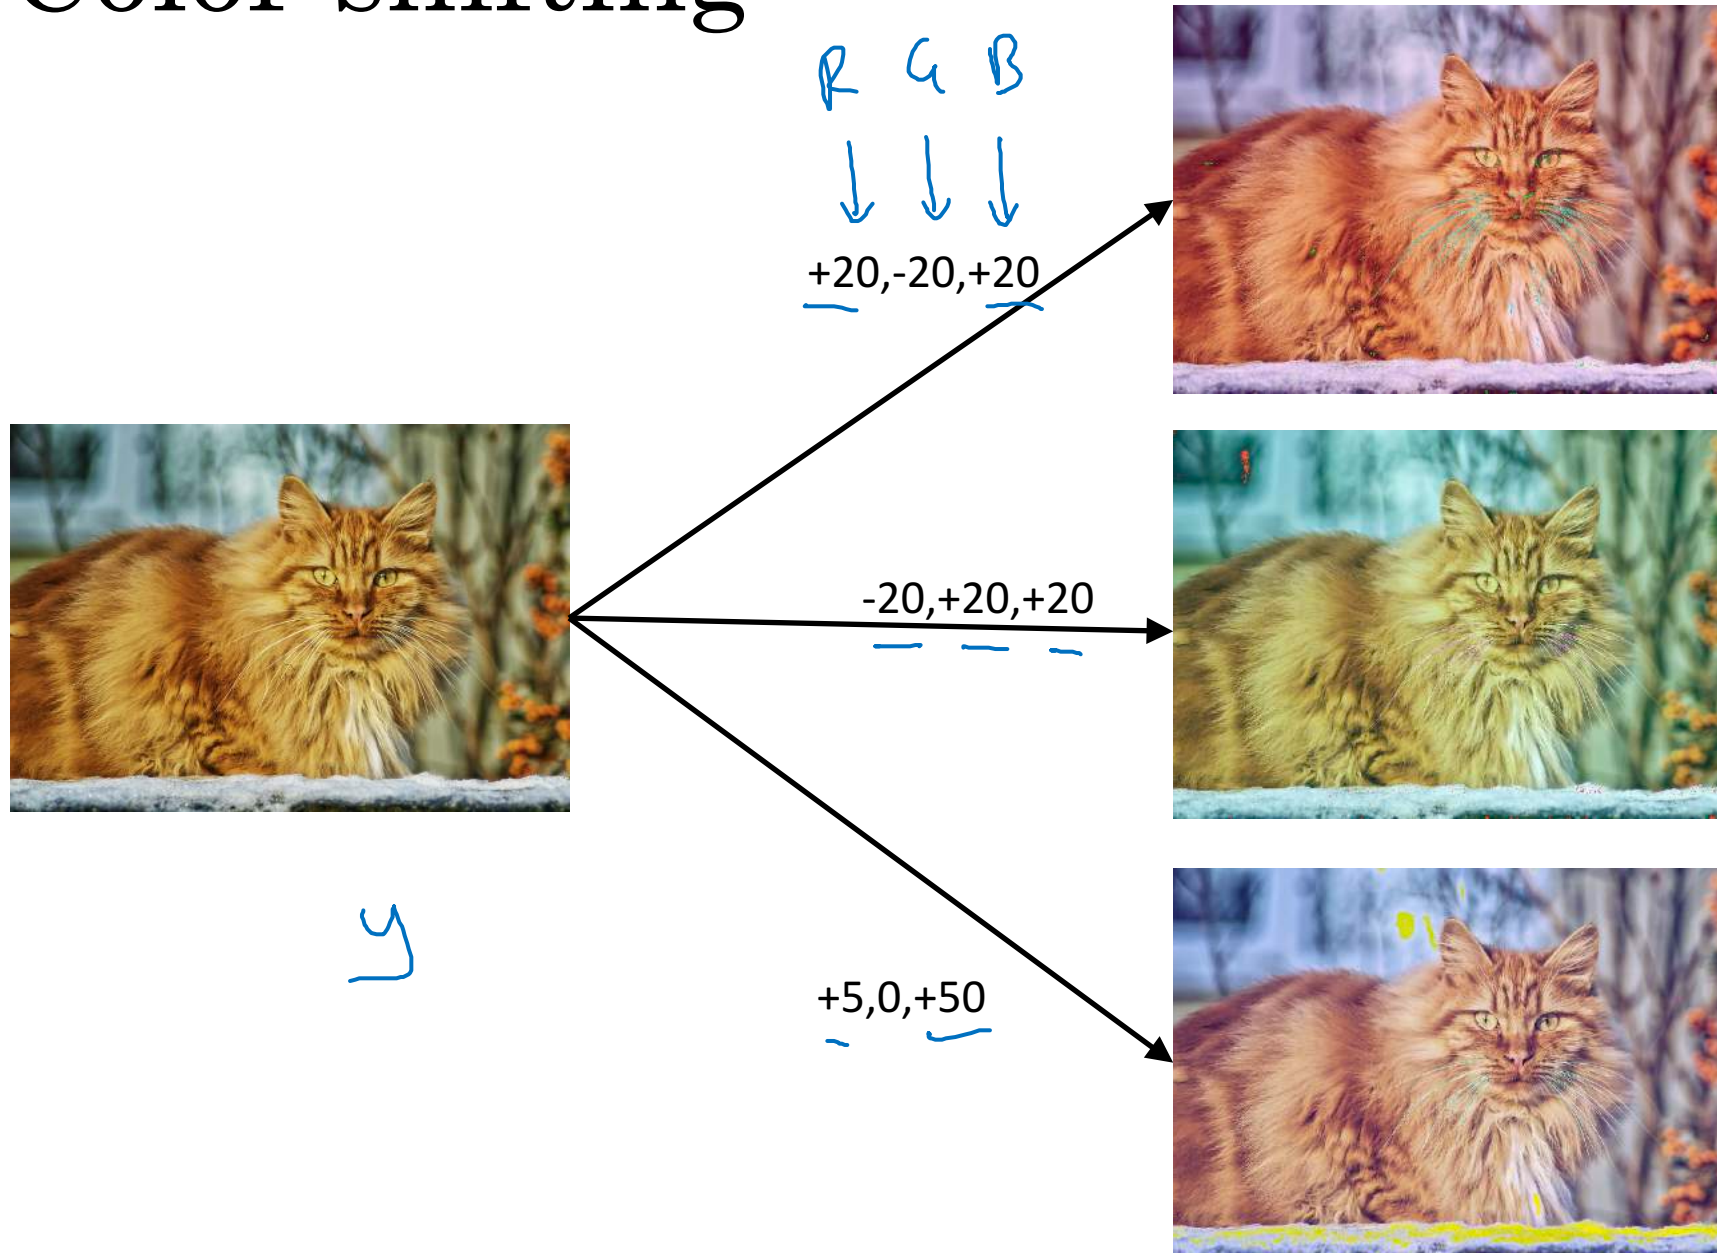
Shearing

Local warping

...



Color shifting



Advanced:

PCA

ml-class.org

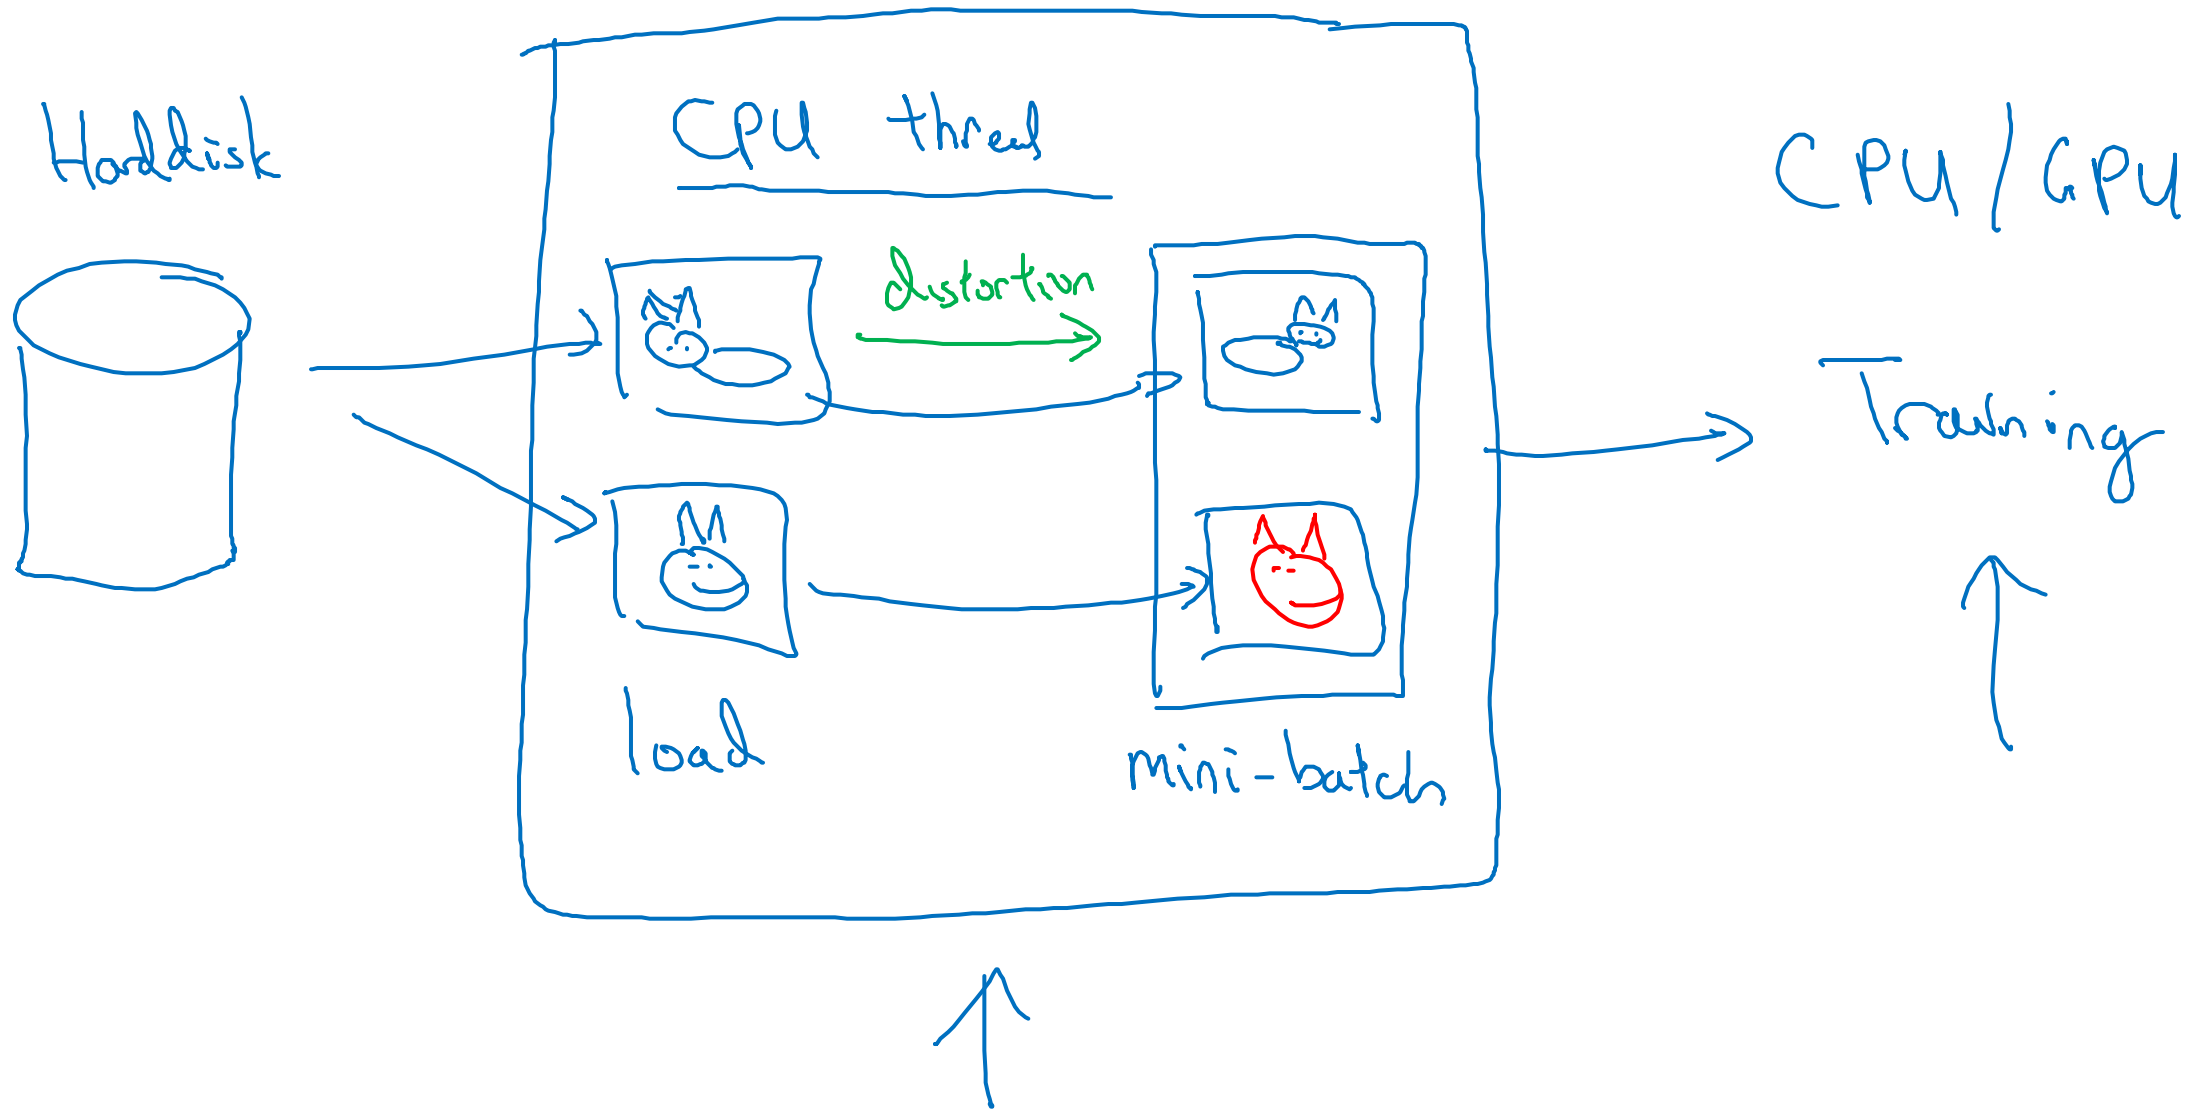
[AlexNet paper

["PCA color augmentation."

R B

G

Implementing distortions during training



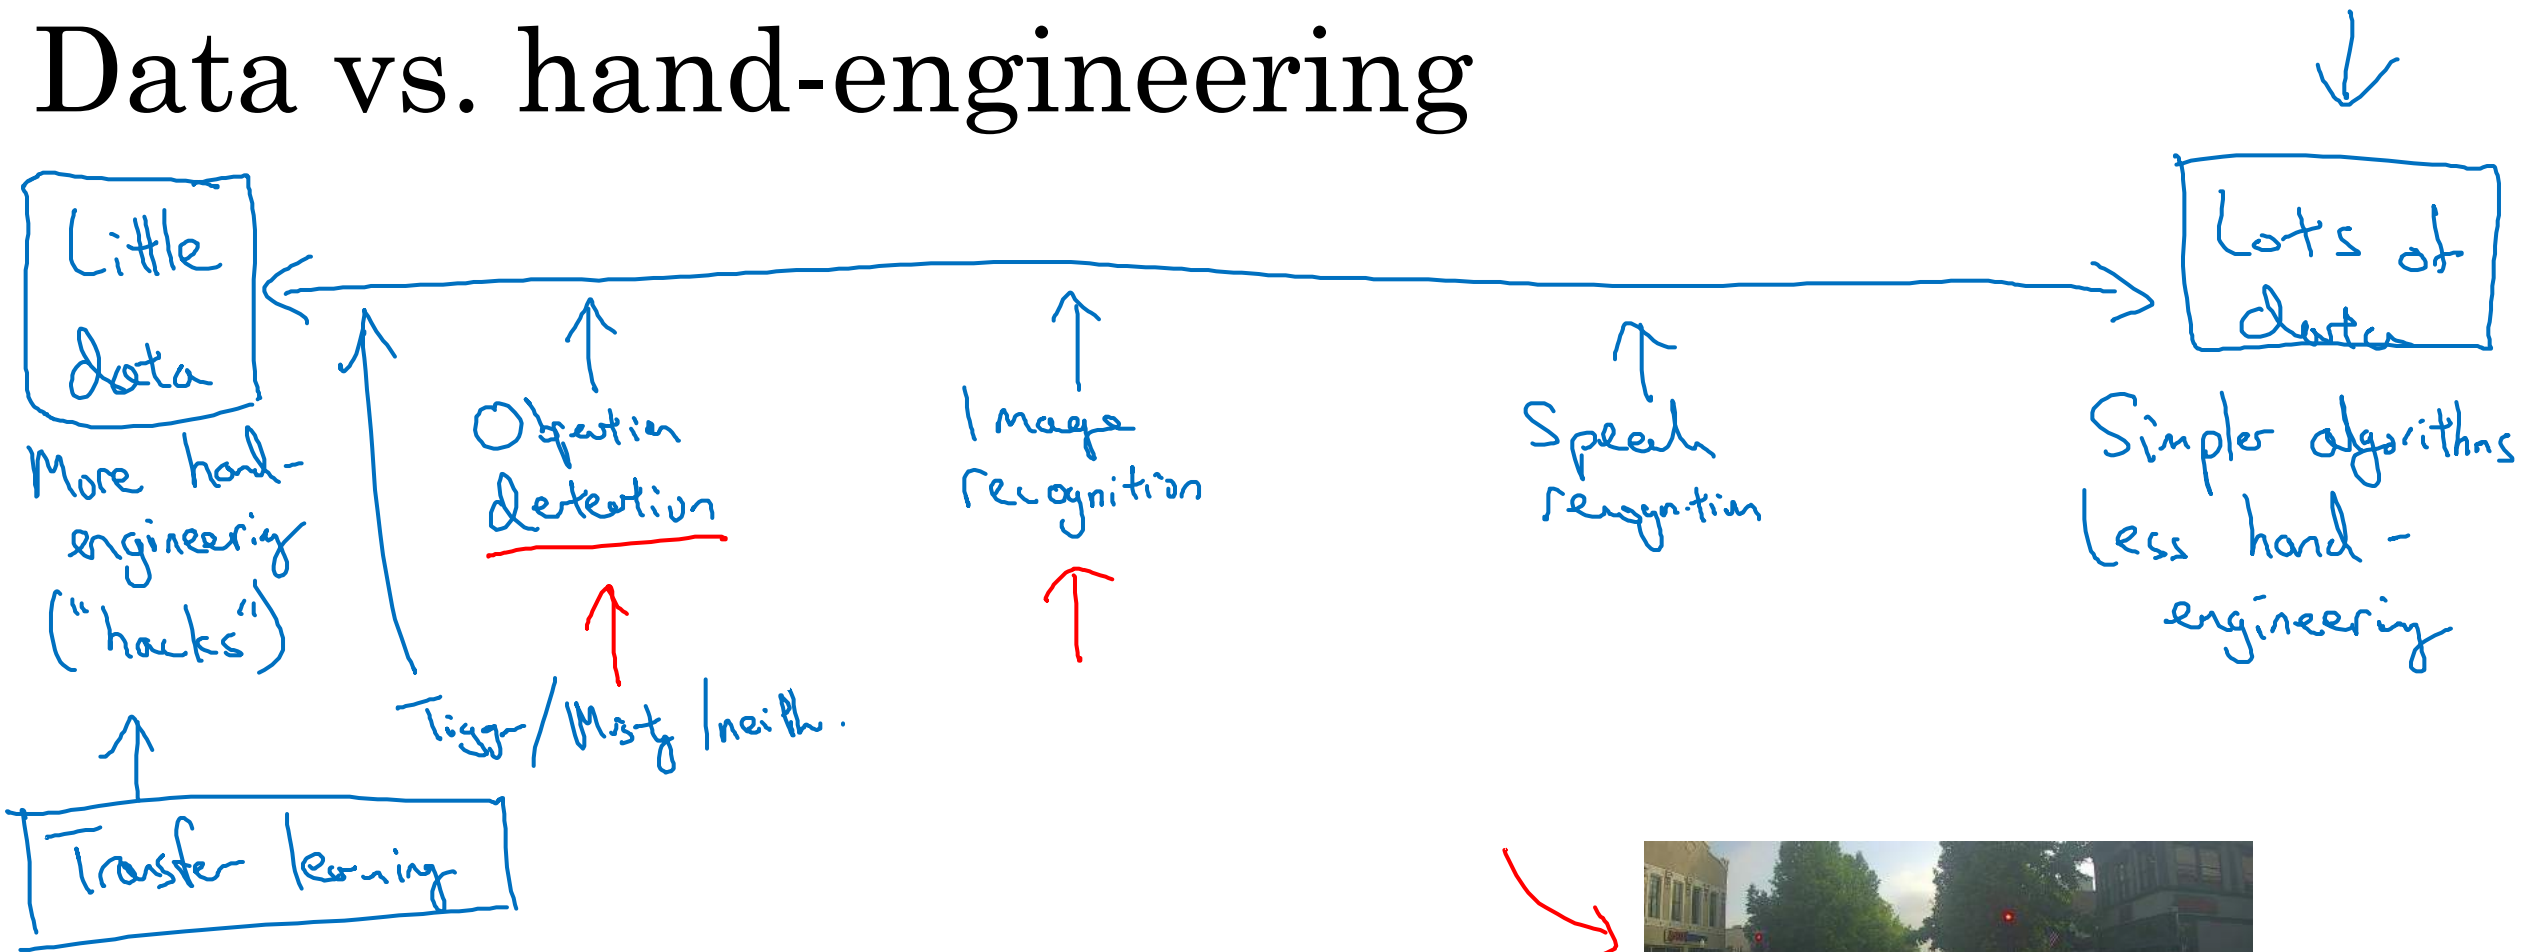


deeplearning.ai

Practical advice for using ConvNets

The state of computer vision

Data vs. hand-engineering



Two sources of knowledge

- • Labeled data (x, y)
- • Hand engineered features/network architecture/other components



Tips for doing well on benchmarks/winning competitions

Ensembling

3-15 networks

→ \hat{y}

- Train several networks independently and average their outputs

Multi-crop at test time

- Run classifier on multiple versions of test images and average results

10-crop



1

+



4

+



1

+



4

Use open source code

- Use architectures of networks published in the literature
- Use open source implementations if possible
- Use pretrained models and fine-tune on your dataset