

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



deeplearning.ai

# Object Detection

---


Object  
localization

# What are localization and detection?

Image classification



Classification with  
localization



Detection



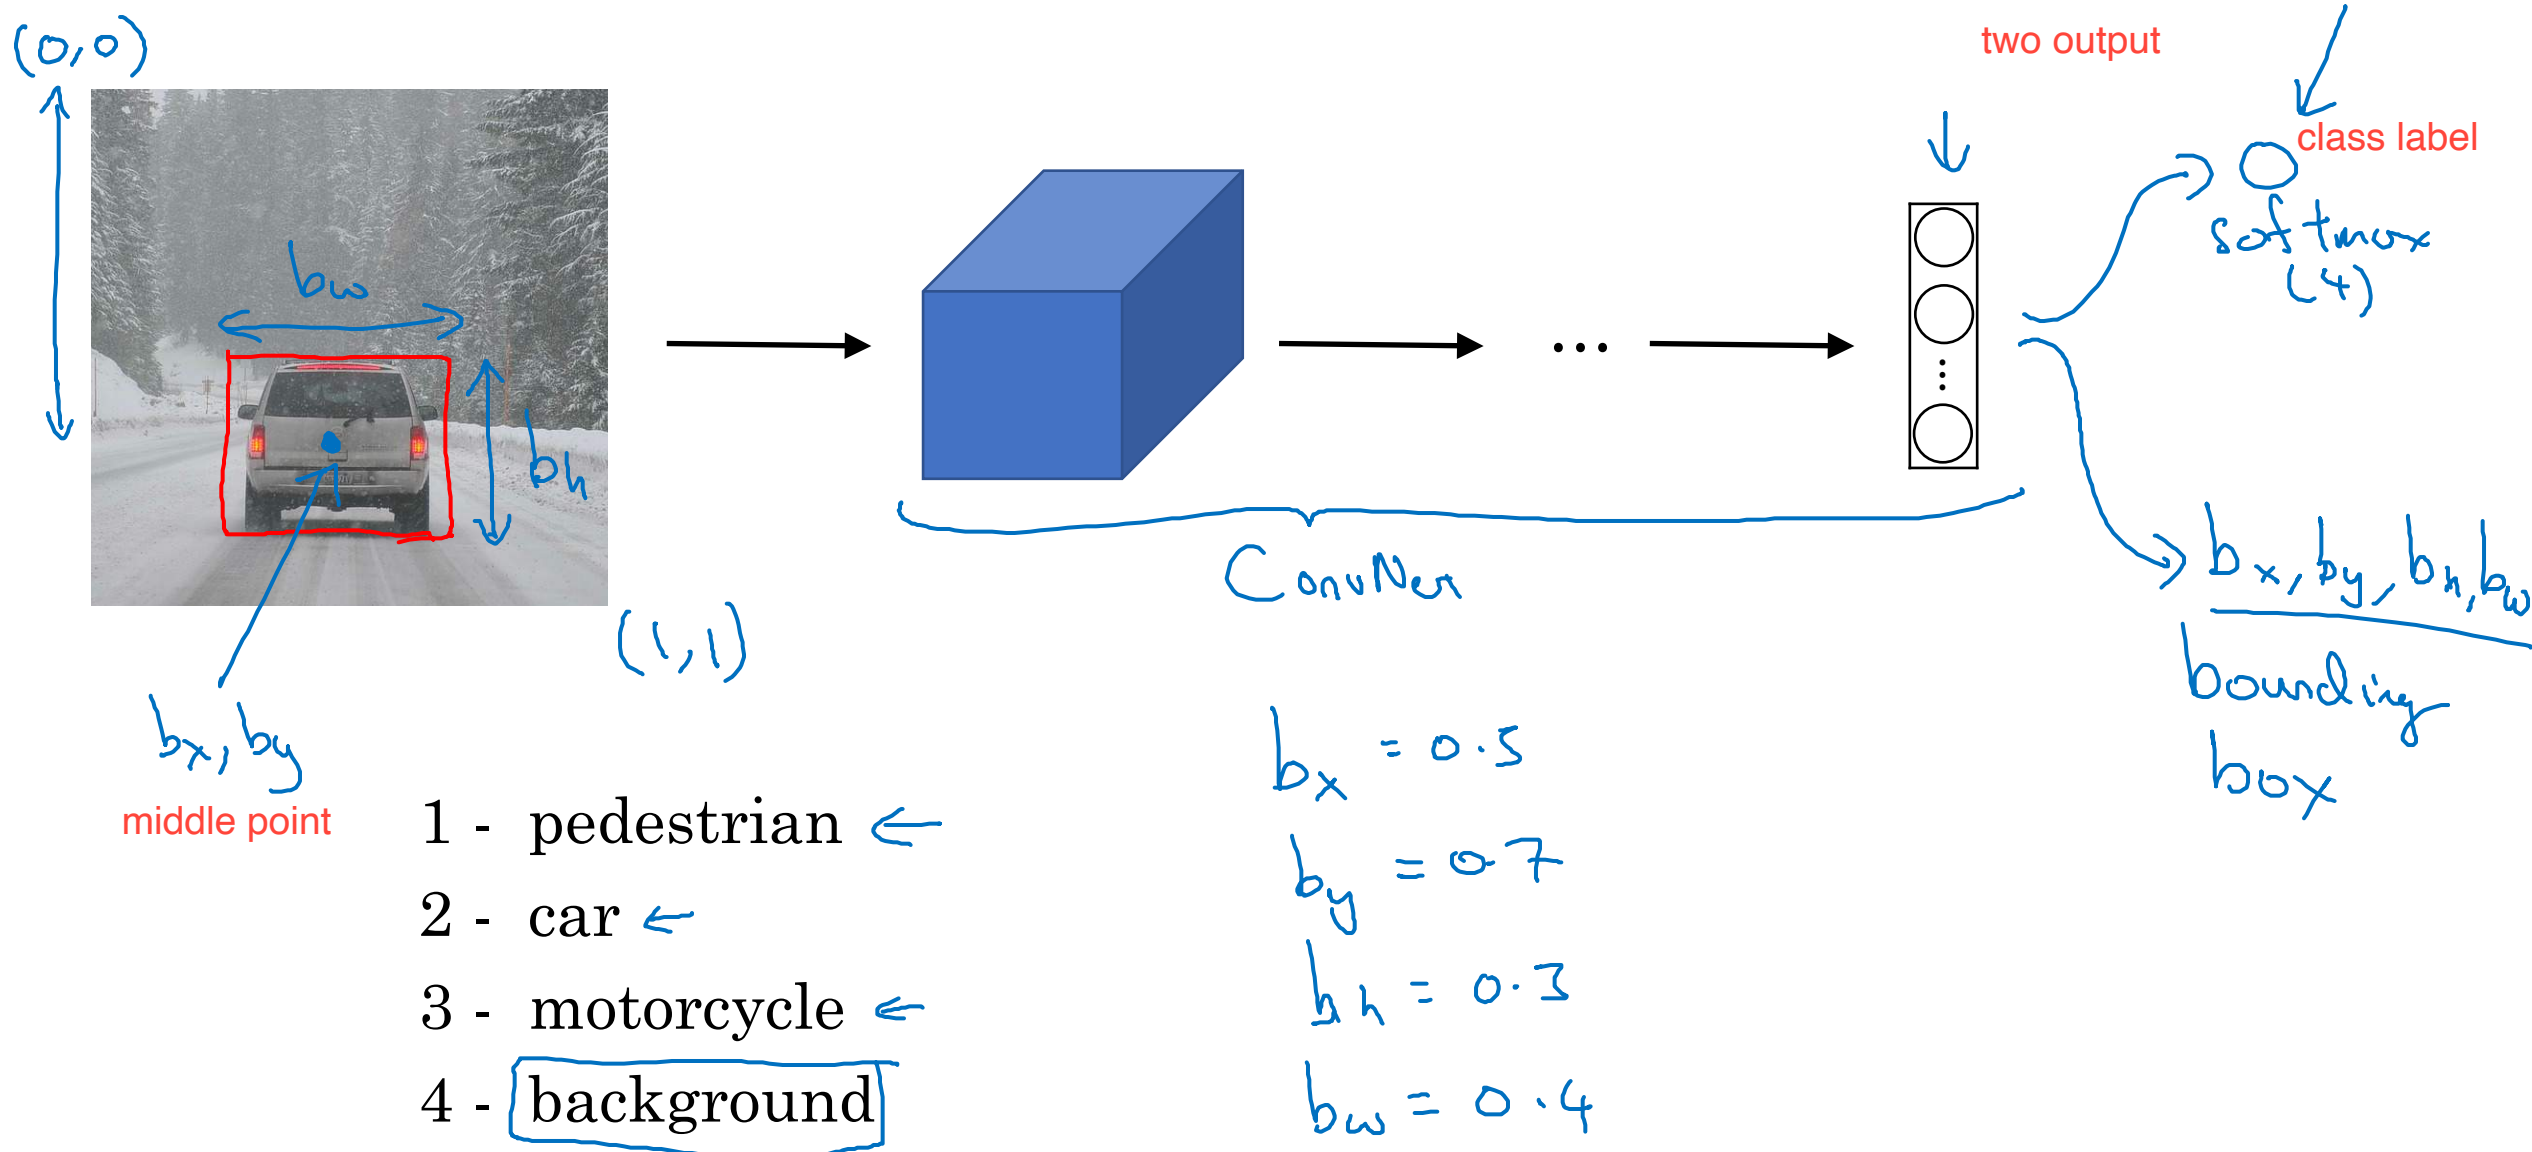
"Car"

"Car"

1 object

multiple  
objects

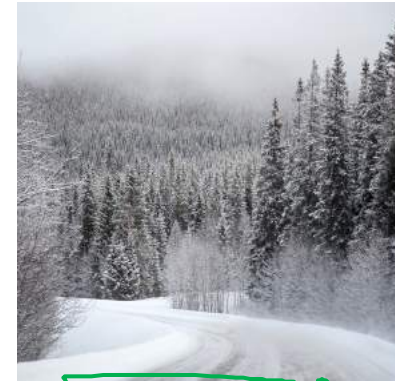
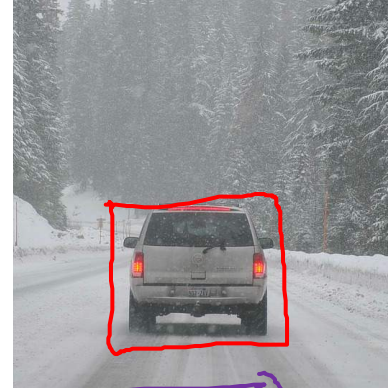
# Classification with localization



# Defining the target label $y$

- 1 - pedestrian
- 2 - car ←
- 3 - motorcycle
- 4 - background ←

Need to output  $b_x, b_y, b_h, b_w$ , class label (1-4)



$$L(\hat{y}, y) = \begin{cases} (\hat{y}_1 - y_1)^2 + (\hat{y}_2 - y_2)^2 + \dots + (\hat{y}_8 - y_8)^2 & \text{if } \underline{y_1 = 1} \\ (\hat{y}_1 - y_1)^2 & \text{if } \underline{y_1 = 0} \end{cases}$$

$$\rightarrow y = \begin{bmatrix} P_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix}$$

is there any object?

training data  
- class labels  
- bounding box

$(x, y)$

$$\begin{bmatrix} 1 \\ b_x \\ b_y \\ b_h \\ b_w \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \end{bmatrix}$$

← "don't care"



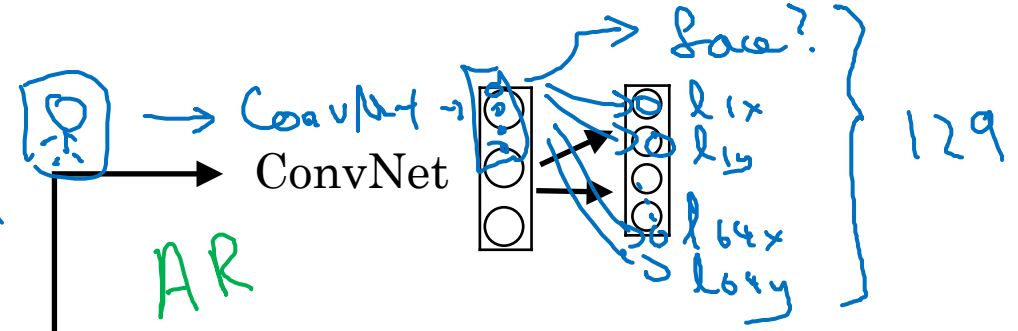
deeplearning.ai

# Object Detection

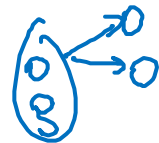
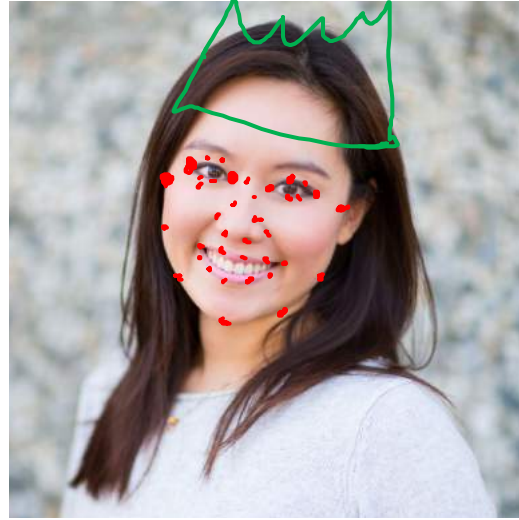
---

Landmark  
detection

# Landmark detection



$b_x, b_y, b_h, b_w$



$l_{1x}, l_{1y}, l_{2x}, l_{2y}, l_{3x}, l_{3y}, l_{4x}, l_{4y}, \dots, l_{64x}, l_{64y}$

$x, y$

$l_{1x}, l_{1y}, \dots, l_{32x}, l_{32y}$





deeplearning.ai

# Object Detection

---

Object  
detection



# Car detection example

Training set:

$X$

$y$



1



1



1



0



0








→ ConvNet →  $y$

# Sliding windows detection

→ ConvNet → 0

→ ConvNet



Computation cost

slow



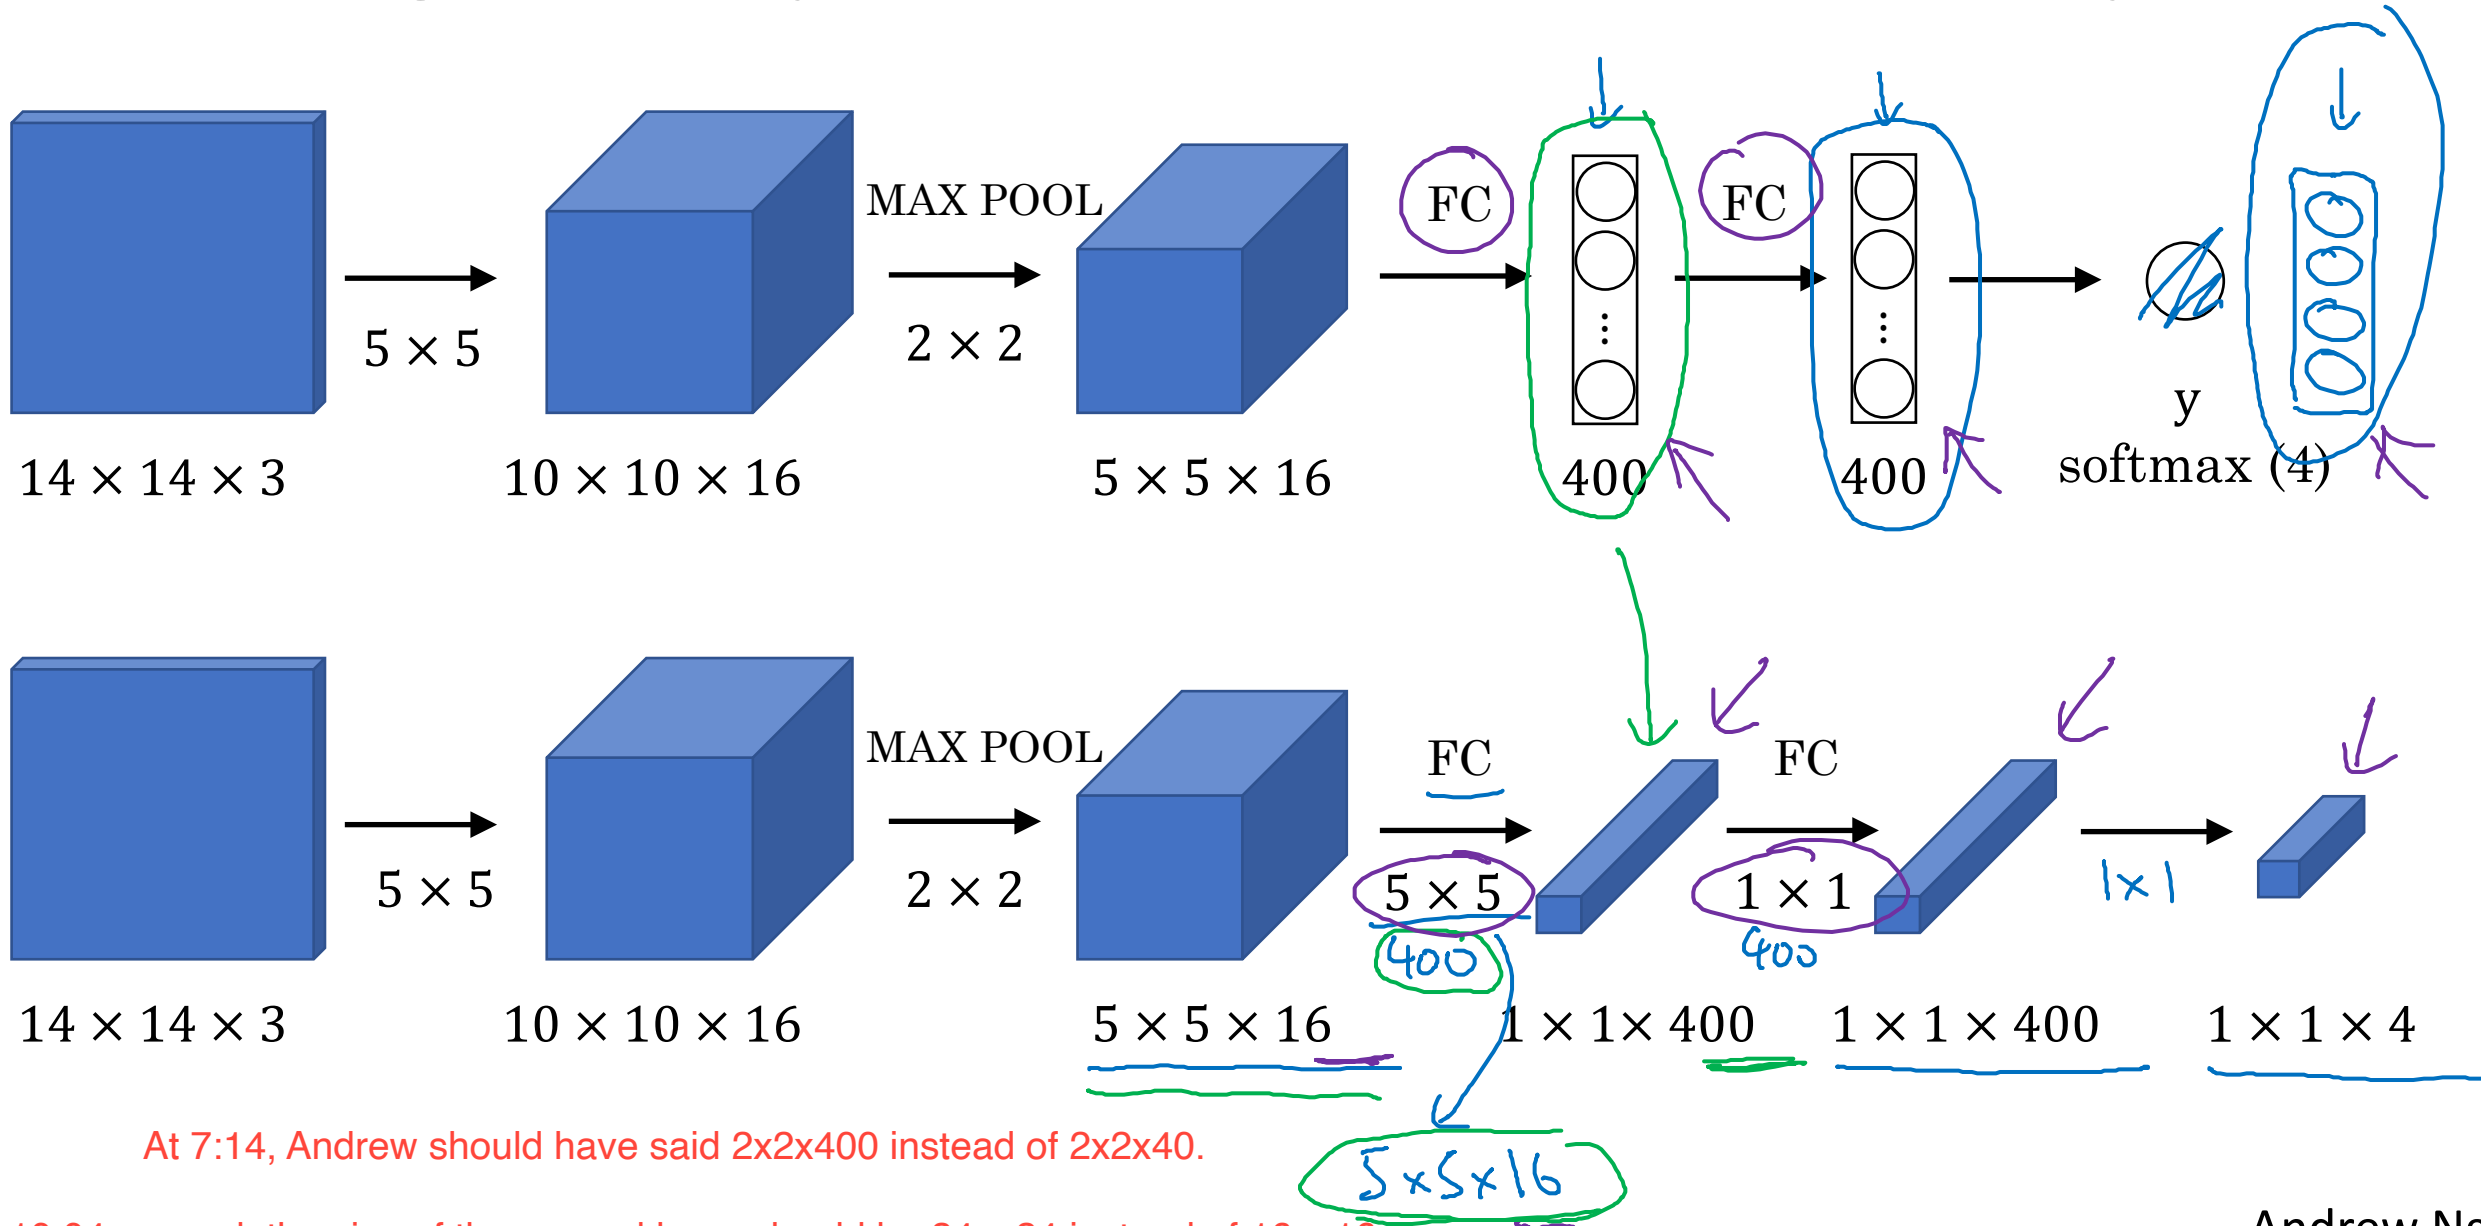
deeplearning.ai

# Object Detection

---

Convolutional  
implementation of  
sliding windows

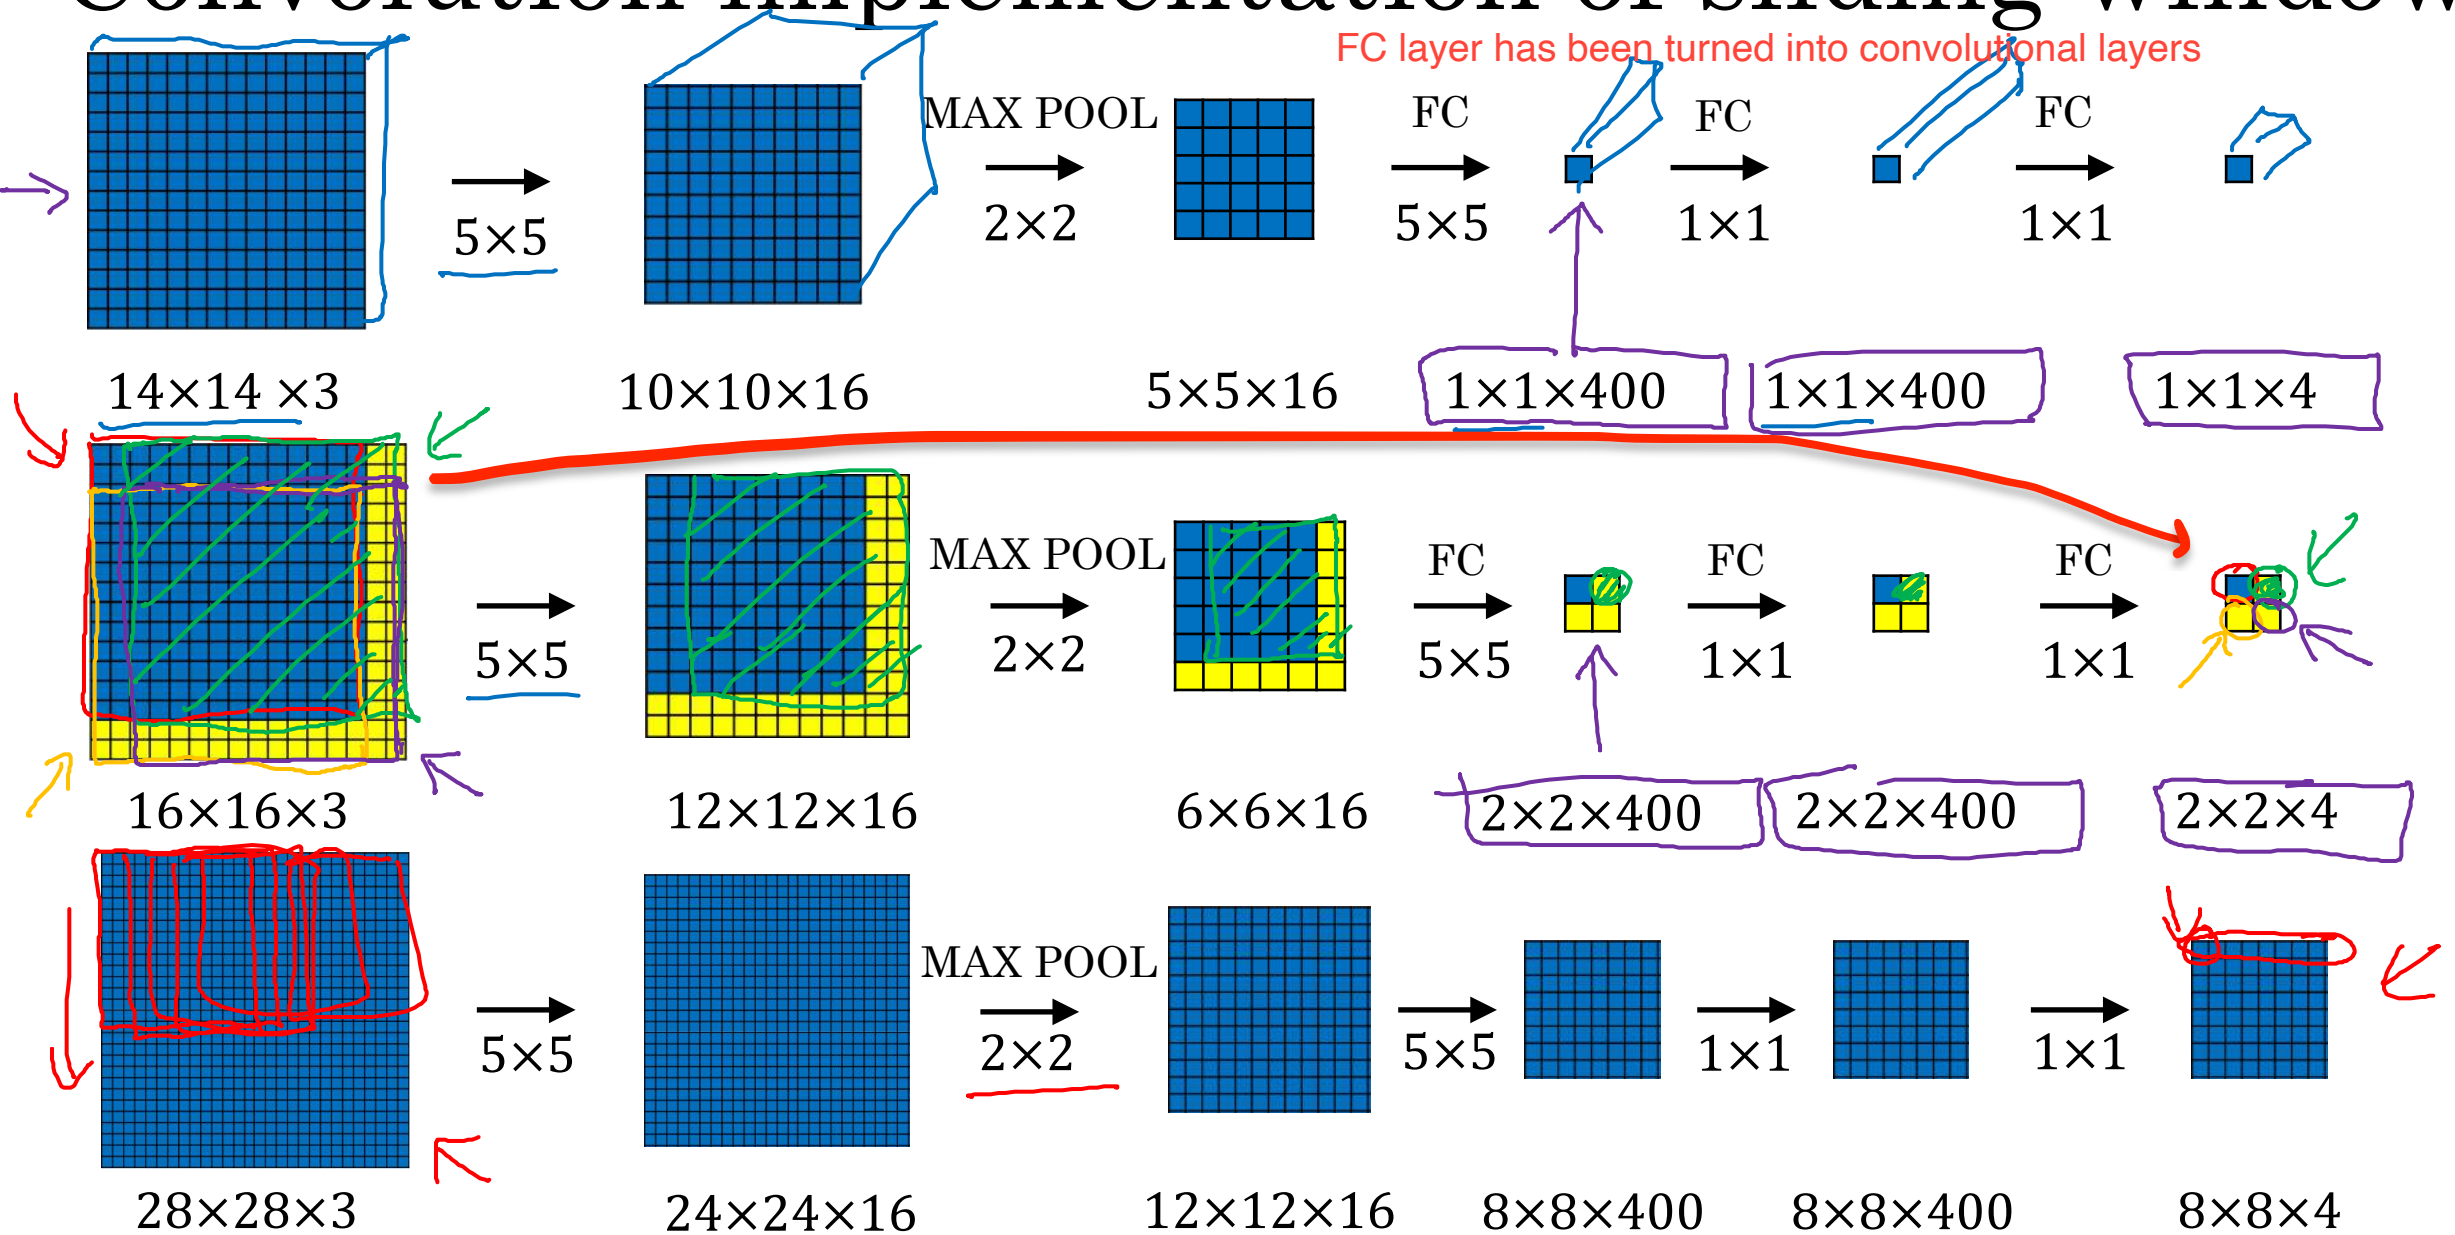
# Turning FC layer into convolutional layers



At 7:14, Andrew should have said  $2 \times 2 \times 400$  instead of  $2 \times 2 \times 40$ .

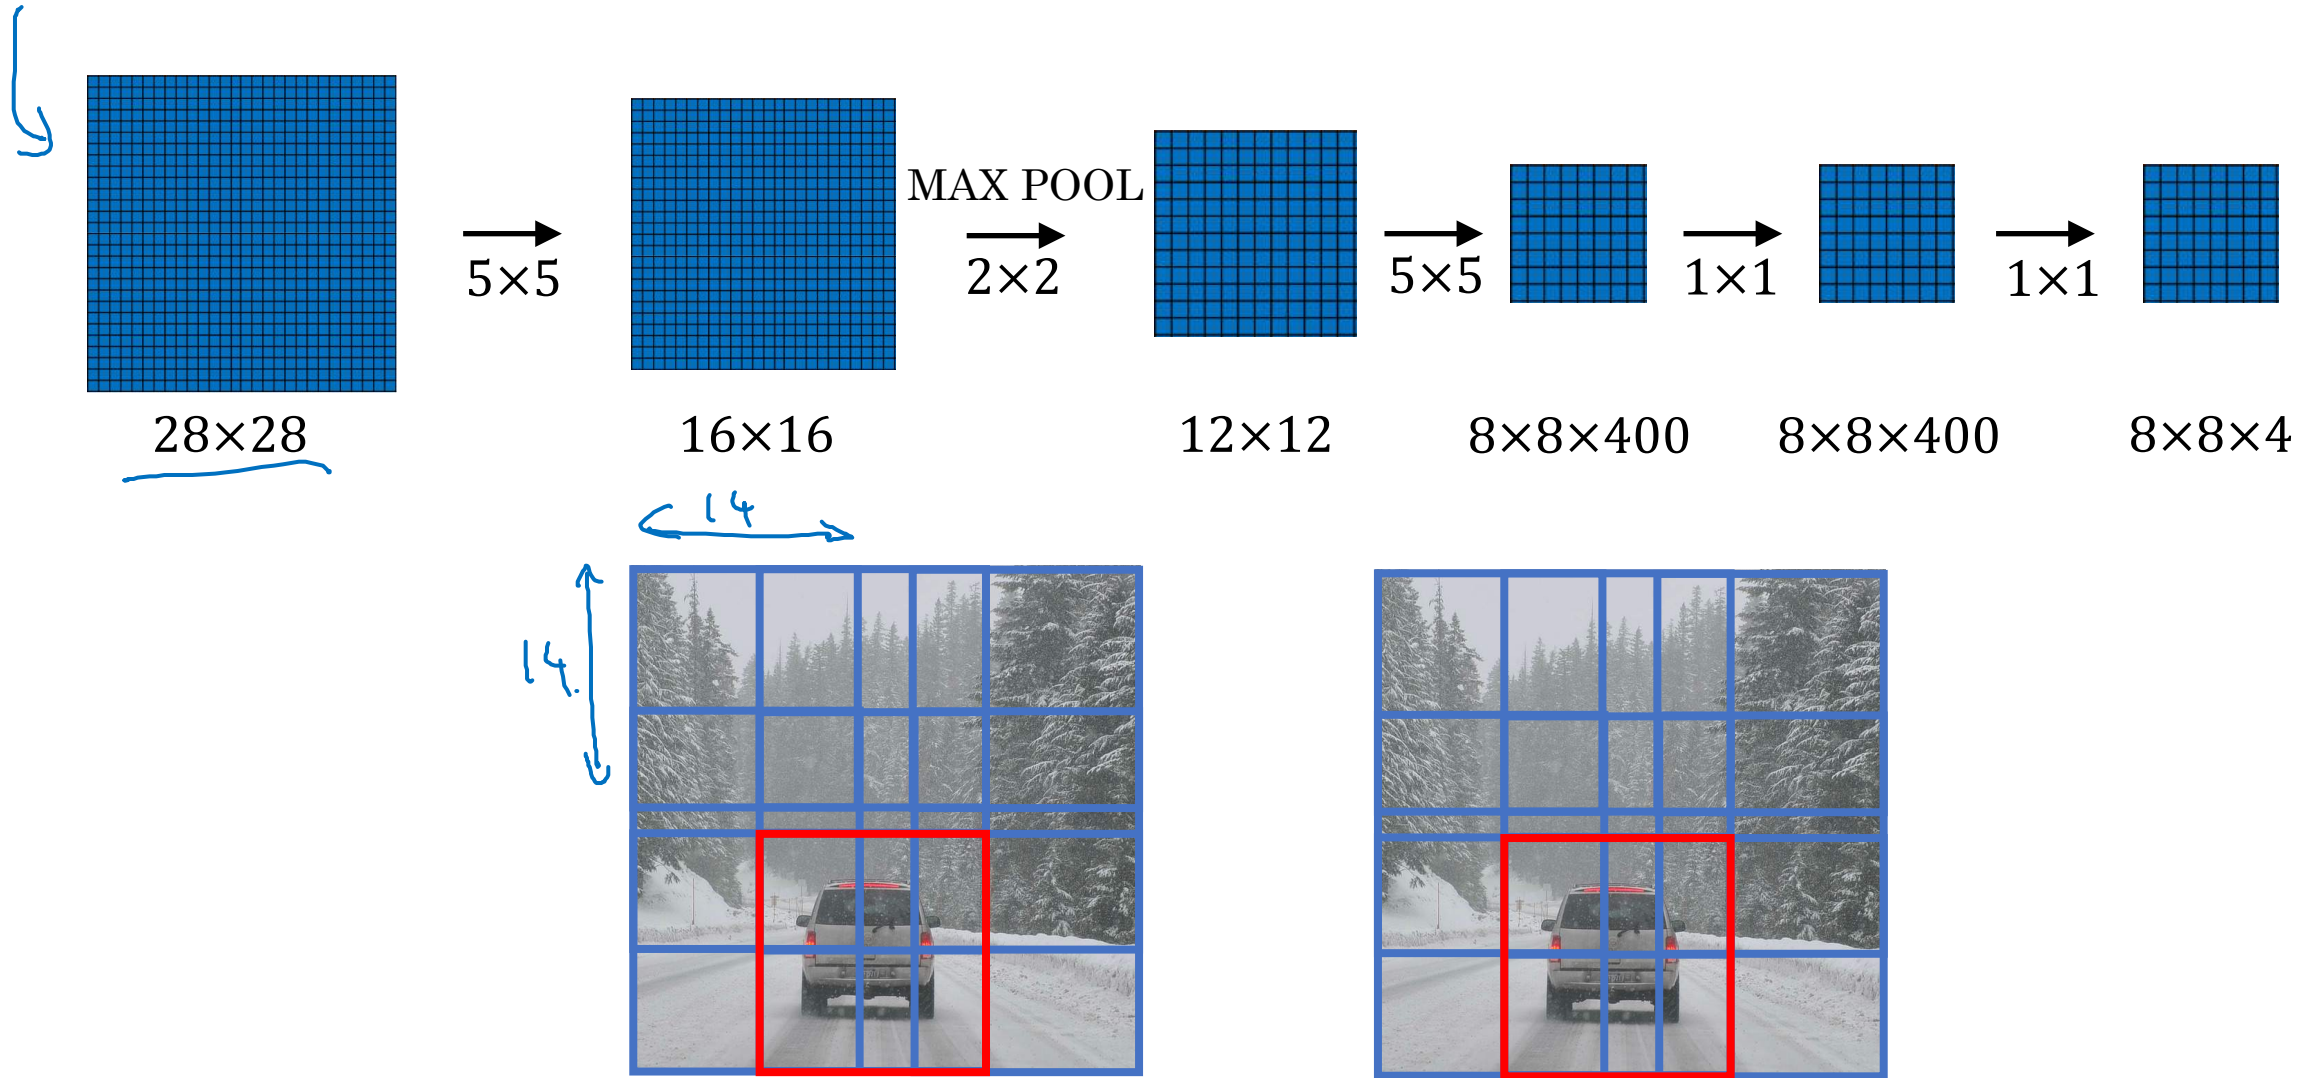
At 10:04 onward, the size of the second layer should be  $24 \times 24$  instead of  $16 \times 16$ .

# Convolution implementation of sliding windows





# Convolution implementation of sliding windows





deeplearning.ai


# Object Detection

---

Bounding box  
predictions




# Output accurate bounding boxes



shape issue  
region issue

# YOLO algorithm



issue: multiple object in a grid cell  
reality: finer grid

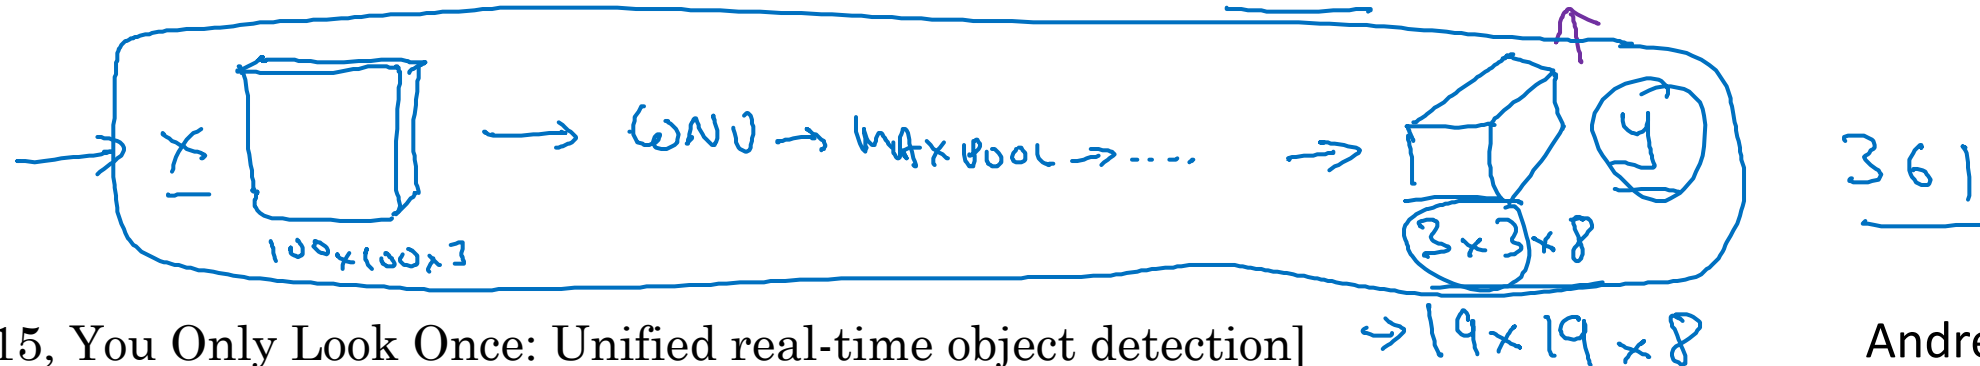


the middle of the bounding box -> use to  
assign to grid

## Labels for training

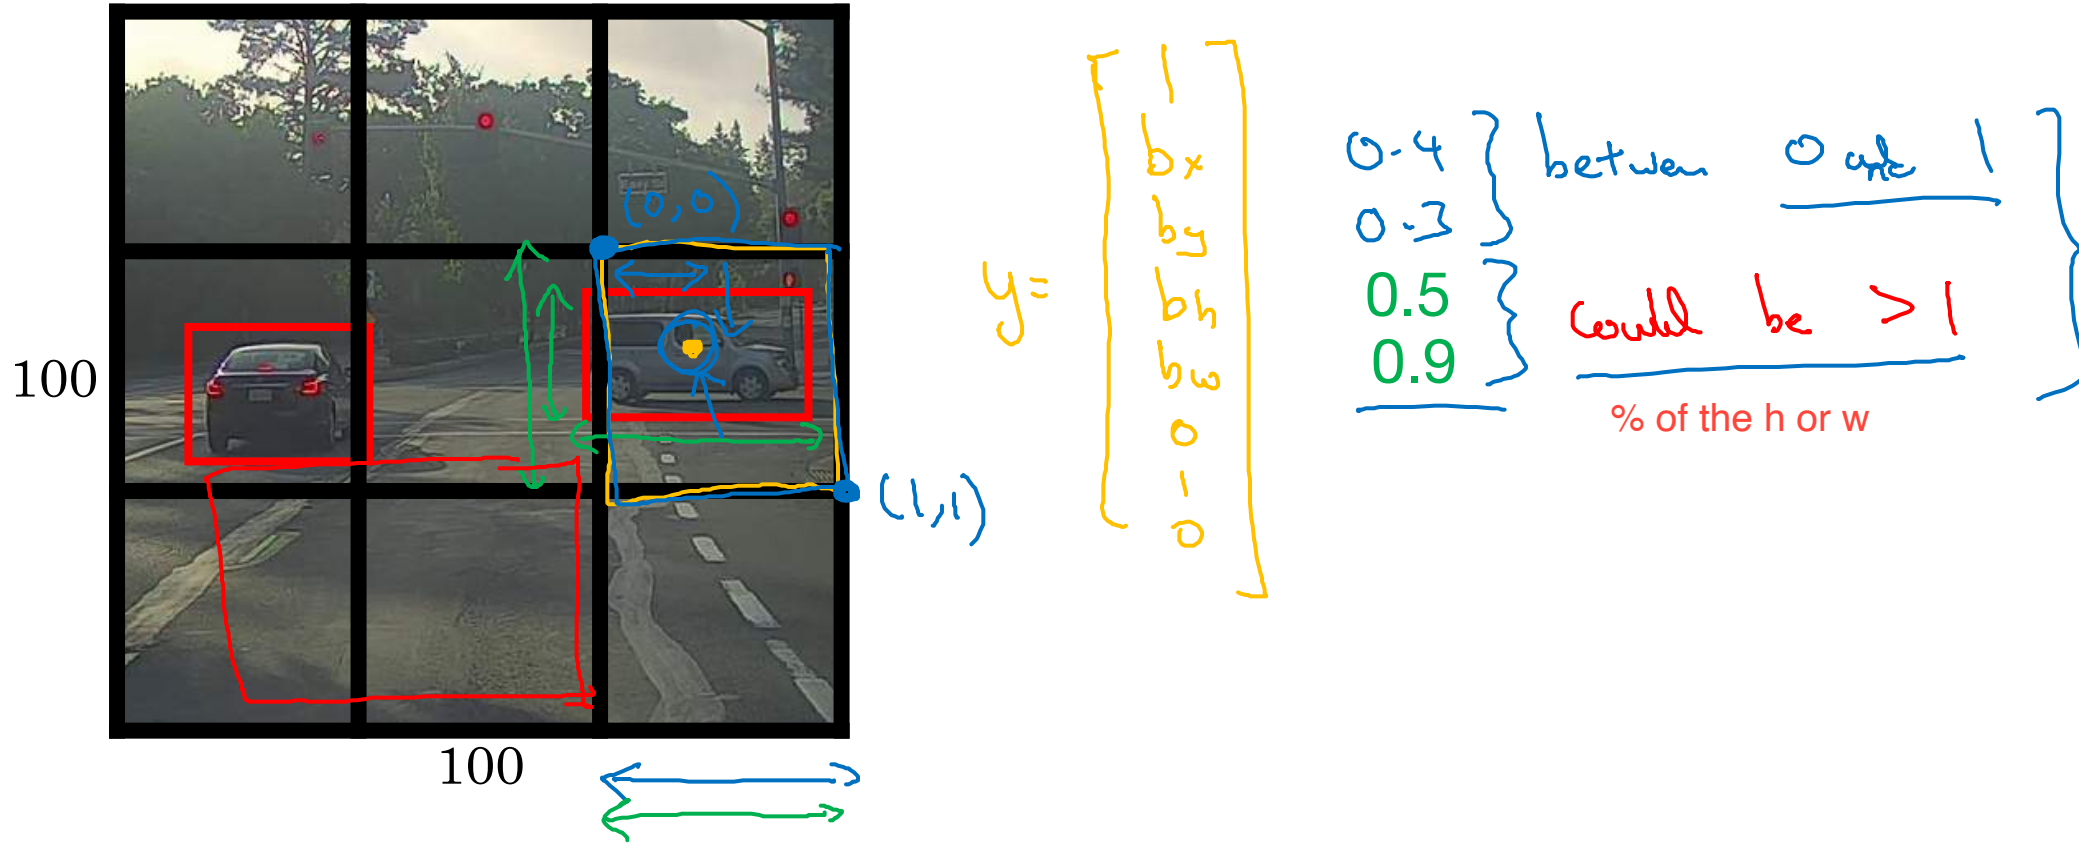
For each grid cell:

Target output:

$3 \times 3 \times 8$



# Specify the bounding boxes





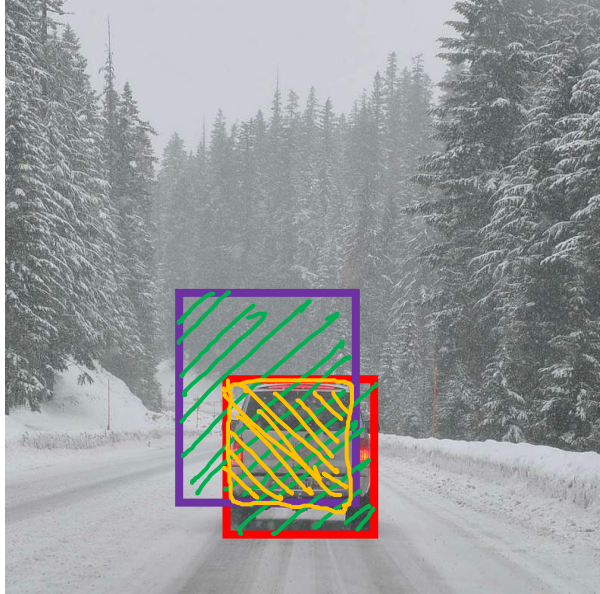
deeplearning.ai

# Object Detection

---

Intersection  
over union

# Evaluating object localization



Intersection over Union (IoU)

$$= \frac{\text{size of } \text{yellow box}}{\text{size of } \text{green box}}$$

“Correct” if  $\text{IoU} \geq 0.5$  ←

“accuracy” for object localization

0.6 ←

More generally, IoU is a measure of the overlap between two bounding boxes.



deeplearning.ai

# Object Detection

---

Non-max  
suppression

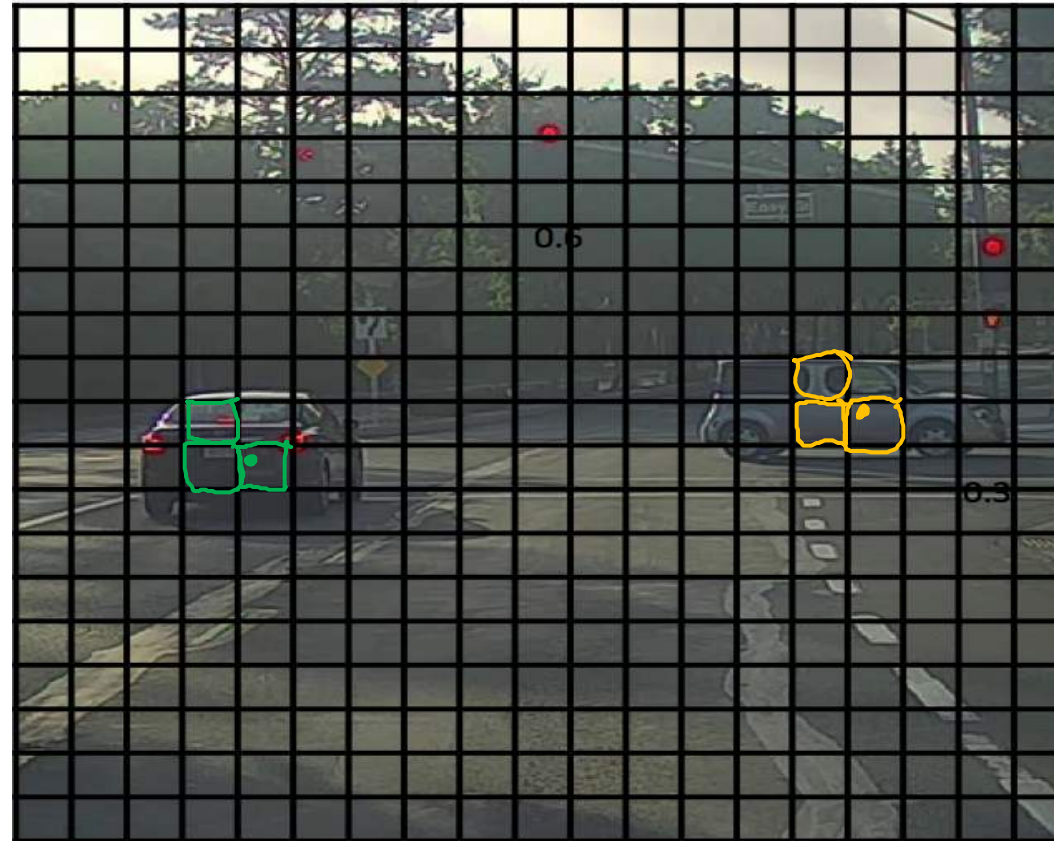


# Non-max suppression example






# Non-max suppression example



19x19


# Non-max suppression example

get just one prediction




$P_c$

# Non-max suppression algorithm



19×19

Each output prediction is:



Discard all boxes with  $p_c \leq 0.6$

→ While there are any remaining boxes:

- Pick the box with the largest  $p_c$   
Output that as a prediction.
- Discard any remaining box with  $\text{IoU} \geq 0.5$  with the box output in the previous step




deeplearning.ai

# Object Detection

---

## Anchor boxes

# Overlapping objects:




$$y = \begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix}$$


Annotations: A green arrow points to  $p_c$ , a blue arrow points to  $b_x$ , and a blue bracket groups  $c_1, c_2, c_3$ .

to handle 2 objects that are in the same grid box


Anchor box 1:



Anchor box 2:



$y =$




# Anchor box algorithm

Previously:

Each object in training image is assigned to grid cell that contains that object's midpoint.

Output  $y$ :  
 $3 \times 3 \times 8$



With two anchor boxes:


Each object in training image is assigned to grid cell that contains object's midpoint and anchor box for the grid cell with highest IoU.

(grid cell, anchor box)



Output  $y$ :  
 $3 \times 3 \times 16$   
 $3 \times 3 \times 2 \times 8$



# Anchor box example



Anchor box 1:      Anchor box 2:



$y =$

$$\begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \\ p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix}$$

1  
bx  
by  
bh  
bw  
1  
0  
0  
1  
bx  
by  
bh  
bw  
0  
0  
0

car only?

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

anchor box 1

anchor box 2





deeplearning.ai


# Object Detection

---


Putting it together:  
YOLO algorithm

# Training

- 1 - pedestrian
- 2 - car ←
- 3 - motorcycle



$y =$



$$\begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \\ p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ 0 \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ 1 \\ b_x \\ b_y \\ b_h \\ b_w \\ 0 \\ 1 \\ 0 \end{bmatrix}$$


$y$  is  $3 \times 3 \times 2 \times 8$

$10 \times 10 \times 16$

$10 \times 10 \times 40$

↑  
#anchors


↑  
5 + #classes



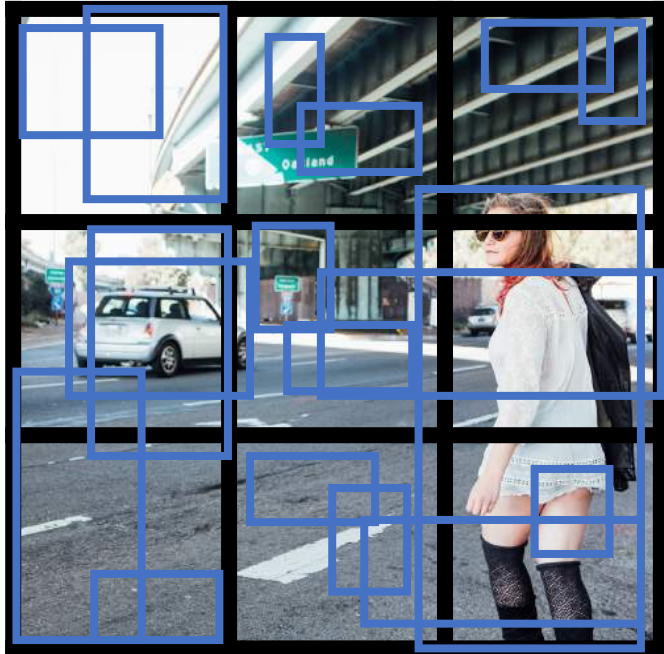
→ ConvNet →



# Making predictions



# Outputting the non-max suppressed outputs



- For each grid cell, get 2 predicted bounding boxes.
- Get rid of low probability predictions.
- For each class (pedestrian, car, motorcycle) use non-max suppression to generate final predictions.



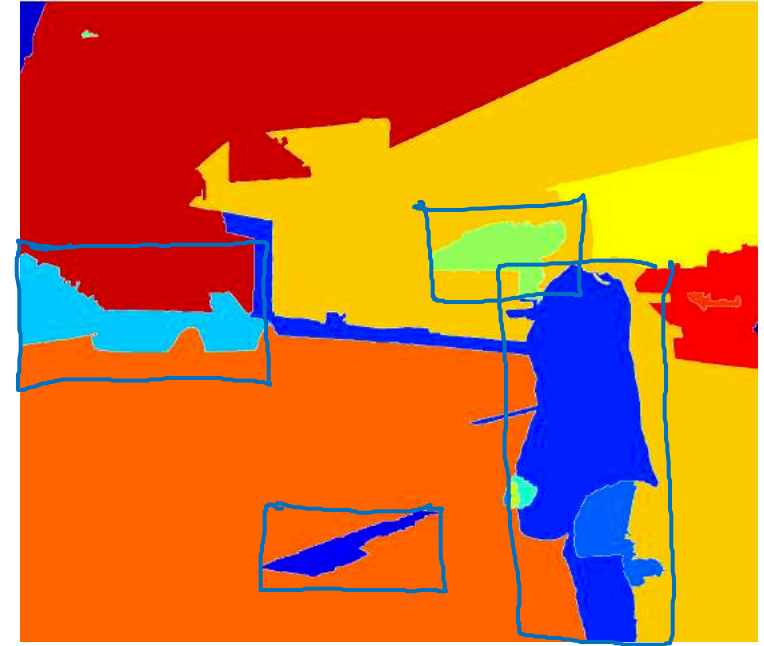
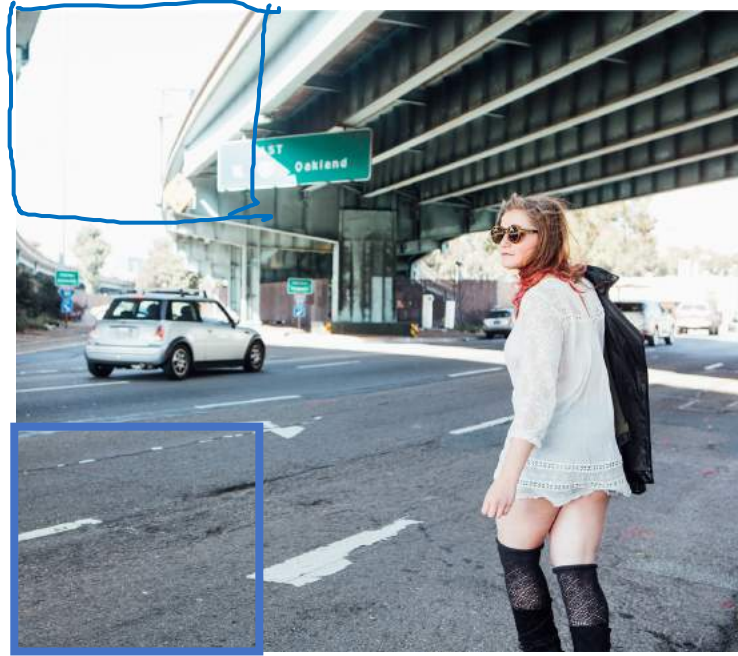
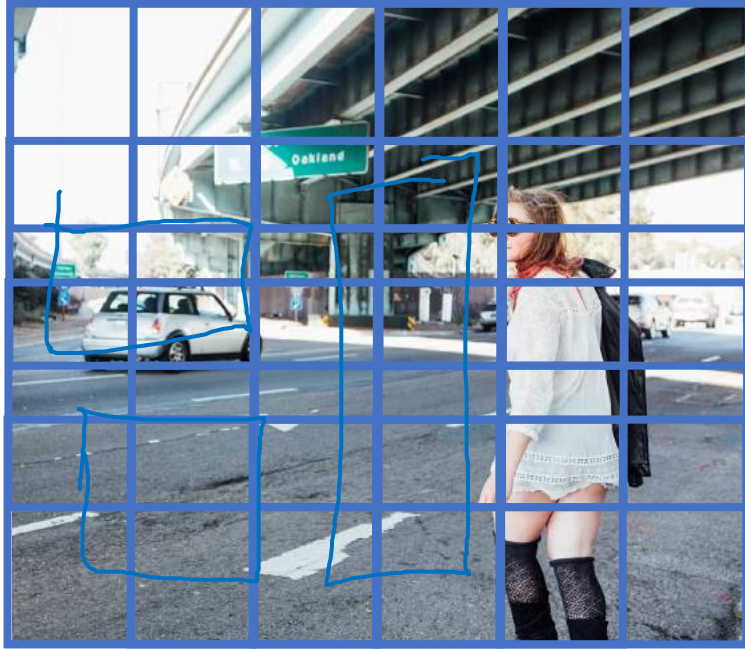
deeplearning.ai

# Object Detection

---

Region proposals  
(Optional)

# Region proposal: R-CNN



Segmentation algorithm

~2,000



# Faster algorithms

segmentation part is slow

→ R-CNN: Propose regions. Classify proposed regions one at a time. Output label + bounding box. ←

Fast R-CNN: Propose regions. Use convolution implementation of sliding windows to classify all the proposed regions. ←

Faster R-CNN: Use convolutional network to propose regions.

slower than yolo

[Girshik et. al, 2013. Rich feature hierarchies for accurate object detection and semantic segmentation]

[Girshik, 2015. Fast R-CNN]

[Ren et. al, 2016. Faster R-CNN: Towards real-time object detection with region proposal networks]

Andrew Ng





deeplearning.ai

# Convolutional Neural Networks

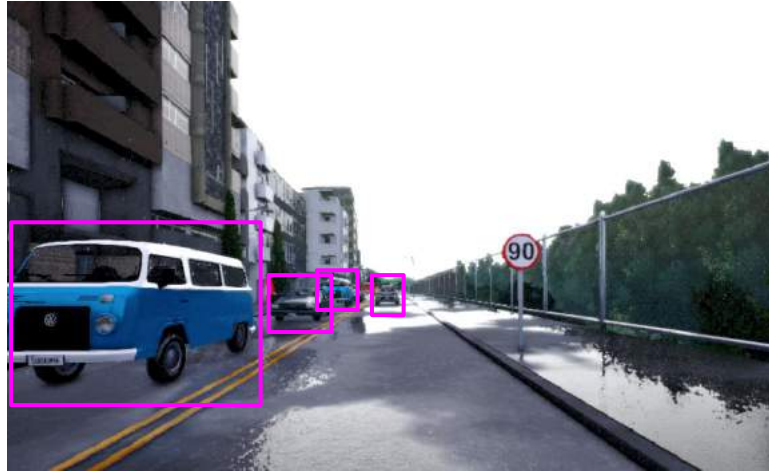
---

Semantic segmentation  
with U-Net


# Object Detection vs. Semantic Segmentation



Input image



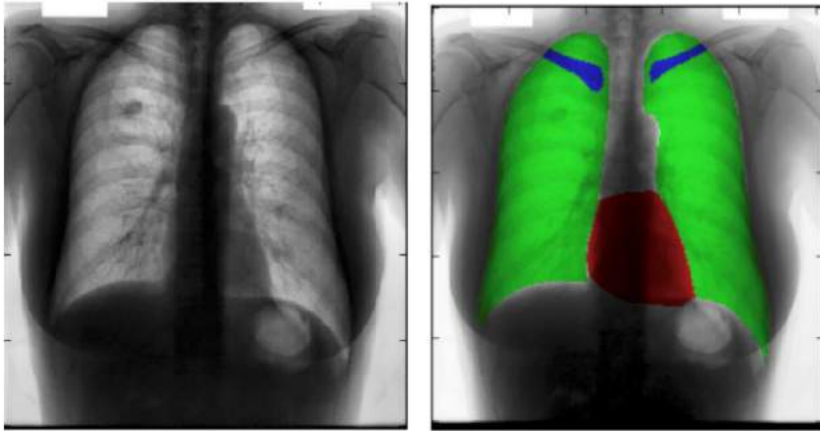
Object Detection



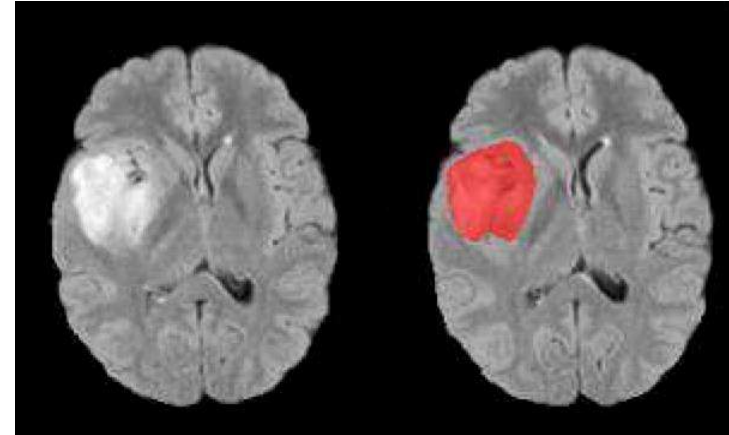
Semantic Segmentation

self driving cars

# Motivation for U-Net



Chest X-Ray

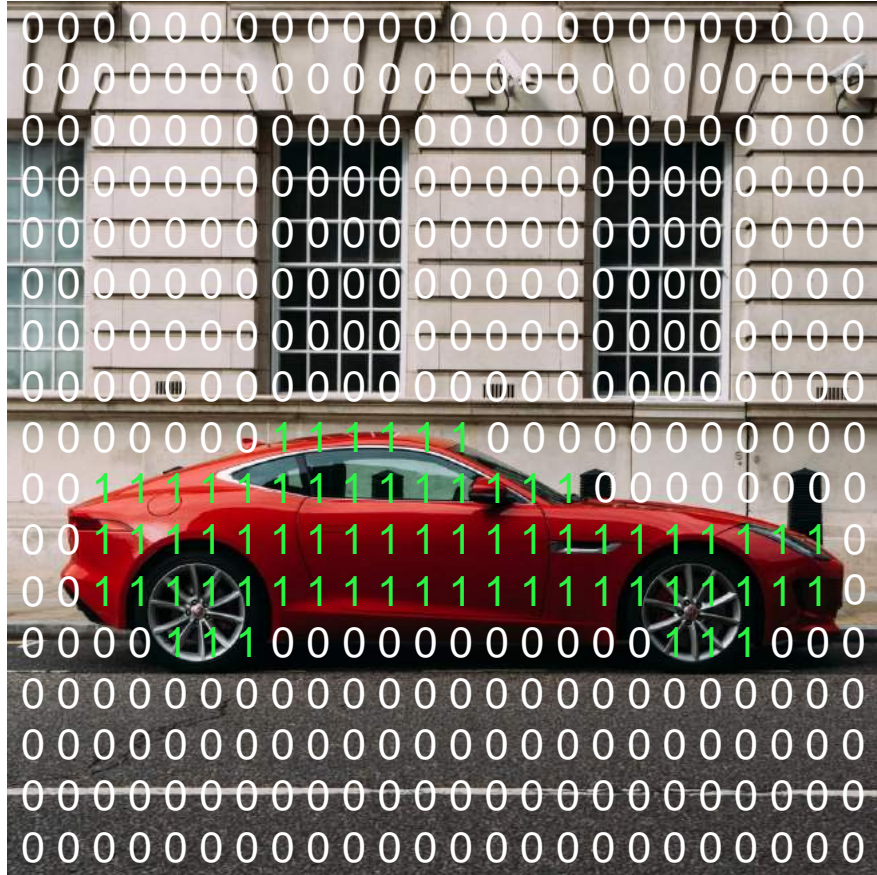


Brain MRI

[Novikov et al., 2017, Fully Convolutional Architectures for Multi-Class Segmentation in Chest Radiographs]


[Dong et al., 2017, Automatic Brain Tumor Detection and Segmentation Using U-Net Based Fully Convolutional Networks ]

# Per-pixel class labels



1. Car  
0. Not Car

# Per-pixel class labels




1. Car
2. Building
3. Road



Segmentation Map



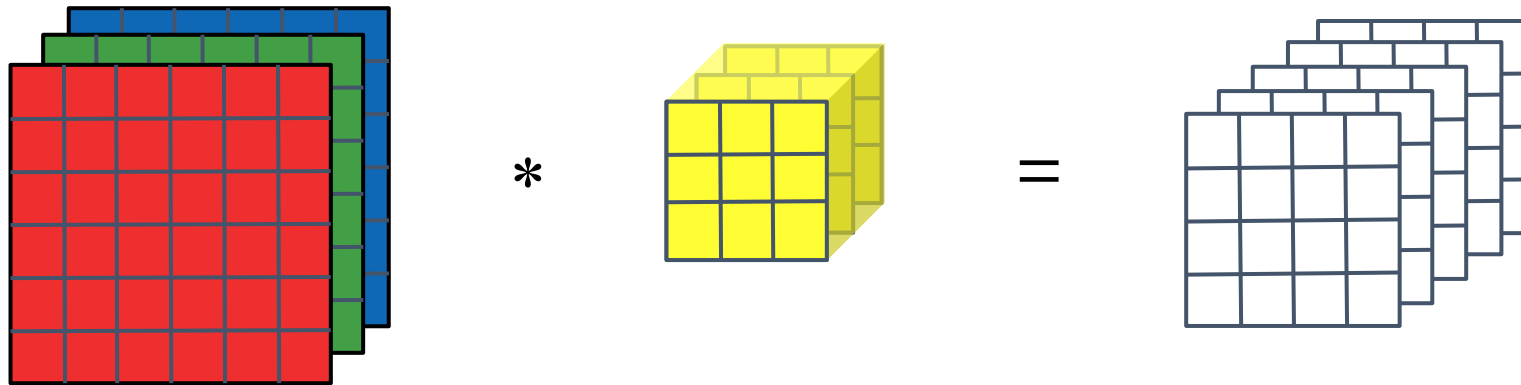
# Deep Learning for Semantic Segmentation



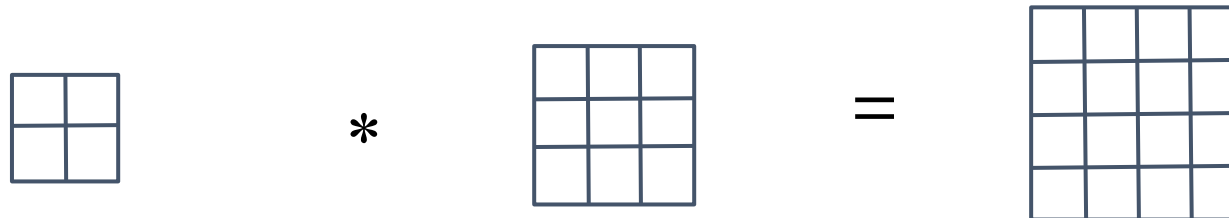


# Transpose Convolution

Normal Convolution



Transpose Convolution



# Transpose Convolution

2	1
3	2


2 x 2

1 <sup>2</sup>	2 <sup>2</sup>	1 <sup>2</sup>
2 <sup>2</sup>	0 <sup>2</sup>	1 <sup>2</sup>
0 <sup>2</sup>	2 <sup>2</sup>	1 <sup>2</sup>

weight filter

ignore padding

add to the overlapping



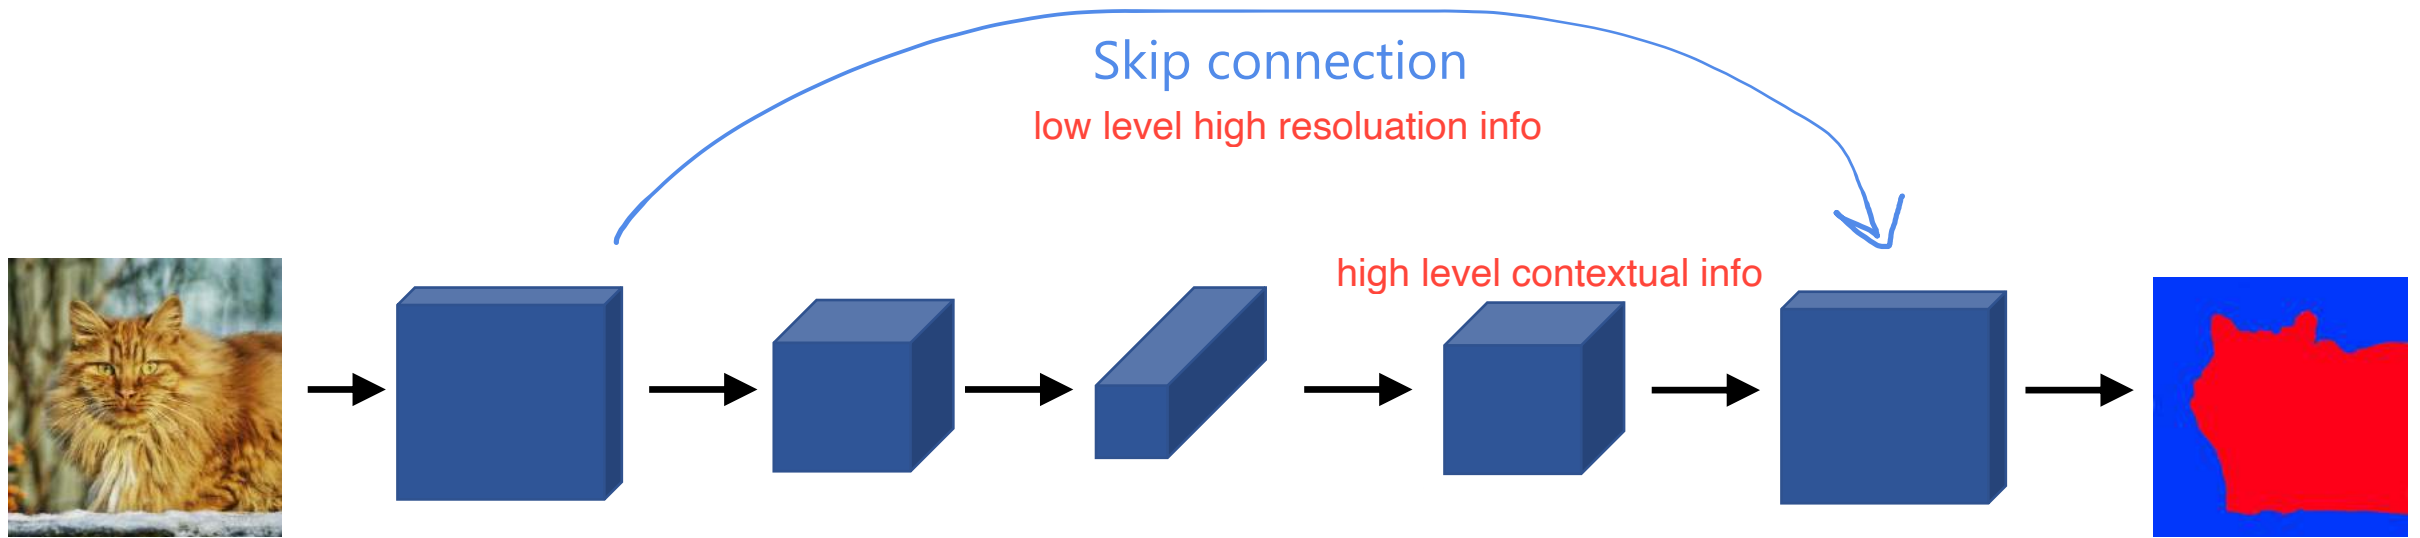
4 x 4

filter  $f \times f = 3 \times 3$


padding  $p = 1$

stride  $s = 2$

# Deep Learning for Semantic Segmentation









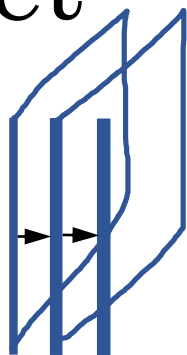
# U-Net



# U-Net



$h \times w \times 3$



$h \times w \times \# \text{ classes}$

- Conv, RELU
- ↓ Max Pool
- ↑ Trans Conv
- Skip Connection
- Conv (1x1)